

# SQL: Part 2



# Simple Aggregation

- ❖ The `count(*)` counts all rows in the table *sp*
- ❖ `count(expression)` counts all non-null values in column

name	item	price
S2	P3	100
S1	P2	20
S1	P1	10
S3	P4	1000
S2	P1	11
S4	P1	9
	P1	
S5		
S4	P3	

```
select count(*) as row_count, count(name) as name_count,  
       count(item) as item_count, count(price) as price_count  
from sp;
```

row_count	name_count	item_count	price_count
9	8	8	6
(1 row)			



---

# Aggregation

---

- ❖ The beginning of data analysis
- ❖ Extremely common (over 90% of queries in TPC-DS)
- ❖ SUM, COUNT, MIN, MAX
- ❖ AVG, STDDEV (samp / pop), VARIANCE, etc.
- ❖ COVAR, CORREL
- ❖ MEDIAN, PERCENTILE, etc.



# GROUP BY

- ❖ FOR EACH unique value of GROUP BY columns
- ❖ *Virtually* segment the table into groups for each unique value
- ❖ Calculate aggregates over those segments

item	price
P1	9
P1	10
P1	?
P1	11
P2	20
P3	?
P3	100
P4	1000
?	?
(9 rows)	

```
select item, avg(price)
from sp
group by item;
```

item	avg
P2	20.0
P1	10.0
P4	1000.0
P3	100.0
(5 rows)	



# GROUP BY

- ❖ one or more columns (or expressions)
- ❖ many systems would let you specify ordinal reference
- ❖ GROUP BY => aggregation
- ❖ Only columns (and expressions based on them) from GROUP BY and aggregate functions can be present in the SELECT list

```
select 'Part ' || item,  
       avg(price)  
from sp  
group by 1;
```

?column?	avg
Part P1	10.0
Part P2	20.0
Part P3	100.0
Part P4	1000.0

(5 rows)



# GROUP BY & DISTINCT

```
select name, item
from sp
group by name, item;
```

name	item
S3	P4
S1	P2
	P1
S1	P1
S2	P3
S2	P1
S4	P1
S5	
S4	P3

(9 rows)

```
select distinct name, item
from sp;
```

name	item
S3	P4
S1	P2
	P1
S1	P1
S2	P3
S2	P1
S4	P1
S5	
S4	P3

(9 rows)

- ❖ SELECT DISTINCT can be done using GROUP BY
- ❖ GROUP BY => aggregation, even without functions



# GROUP BY & HAVING

- ❖ HAVING clause is a way to specify “WHERE” to results of aggregation
- ❖ Find items *where* average price is  $> 20$
- ❖ HAVING requires a GROUP BY clause
- ❖ HAVING on GROUP BY columns?

```
select item, avg(price)
from sp
group by item;
```

item	avg
P2	20.0
P1	10.0
P4	1000.0
P3	100.0

(5 rows)

```
select item, avg(price)
from sp
group by item
having avg(price) > 20;
```

item	avg
P4	1000.0
P3	100.0

(2 rows)



---

# GROUP BY & HAVING

---

- ❖ HAVING can have subqueries
- ❖ Find items *where* avg price is  $>$  avg price of all items

```
select item, avg(price)
from sp
group by item
having avg(price) > (select avg(price)
                    from sp)
```

```
;
```

item	avg
P4	1000.000000000000000000

(1 row)



# Rows with Aggregate Property

- ❖ Want to determine item with Minimum Price (not for each item, find the minimum price)
- ❖ Find the row(s), thus item(s), where price = minimum price
- ❖ Is there another way?

```
select item, min(price)
from sp;
ERROR: column "sp.item" must appear
in the GROUP BY clause
or be used in an aggregate function
LINE 1: select item, min(price) from sp;
```

item	price
P1	9
P1	10
P1	?
P1	11
P2	20
P3	?
P3	100
P4	1000
?	?

(9 rows)

```
select item, price
from sp
where price = (select min(price)
               from sp);
```

item	price
P1	9

(1 row)



# Rows with Aggregate Property

- ❖ Find the row(s), thus items(s), where price = minimum price - not using the subquery

```
select item, price
from sp, (select min(price)
          from sp) mp(min_price)
where price = min_price;
```

item	price
P1	9

(1 row)

- ❖ Two are equivalent

- ❖ One more way - using Window functions

```
select item, price
from (select item, price,
             rank() over (order by price) as prank
      from sp) ranked_price
```

```
where prank = 1;
```

item	price
P1	9

(1 row)



# Aggregates with DISTINCT

- ❖ Single argument aggregates (COUNT, SUM, STDDEV\_POP) etc. can take a DISTINCT option
- ❖ COUNT(DISTINCT ...) is the biggest use - to get an idea for number of unique values in columns
- ❖ Conceptually, first compute the DISTINCT values and then do the aggregate

name	item	price
S2	P3	100
S1	P2	20
S1	P1	10
S3	P4	1000
S2	P1	11
S4	P1	9
	P1	
S5		
S4	P3	

```
select count(distinct name),  
       count(distinct item)
```

```
from sp;
```

count	count
5	4
(1 row)	



---

# Available Functions

---

- ❖ Single Argument (Algebraic)
  - ❖ SUM, COUNT, MIN, MAX, AVG, STDDEV\_(POP/SAMP), VAR\_(POP/SAMP), etc.
- ❖ Two Argument
  - ❖ CORREL, COVAR\_(POP/SAMP), REGR\_ family
- ❖ Single Argument (Holistic)
  - ❖ PERCENTILE\_(CONT/DISC), MEDIAN, QUARTILE
- ❖ ANY(boolean expression); EVERY (boolean expression)
  - ❖ T if True for any row in group; T if True for every row in group



---

# Review

---

SELECT [options] column\_expression\_list  
FROM table\_expression\_list  
WHERE condition  
GROUP BY groupby\_list  
HAVING condition

- ❖ Conceptually, we first evaluate the cross-products, joins and WHERE conditions
- ❖ Then we aggregate the rows according to the GROUP BY expression computing functions mentioned in SELECT list and HAVING condition
- ❖ Then we apply the HAVING condition to the resulting rows
- ❖ Finally we evaluate the SELECT list column expressions



---

# Review: *A Few Points*

---

- ❖ The SELECT list can only contain references to GROUP BY list of column\_expressions and aggregate functions
- ❖ Nothing else makes sense
  - ❖ Remember: for each [group by column] compute aggregate
- ❖ No nesting AVG(COUNT(...)) etc. as it doesn't make sense - use derived tables (nested structures)



---

# Derived Tables

---

- ❖ In the FROM clause any query expression can be considered a table
- ❖ Syntax:
  - ❖ ( <query expression> ) AS <dtname> [(column\_name\_list)]
- ❖ Requirements:
  - ❖ All columns must have unique name
    - ❖ via “AS” renaming in SELECT of <query expression>, or
    - ❖ by listing out column names in the derived table name



# Derived Tables - Example 1

- ❖ Find suppliers who offer a price lower or equal to avg price for item

```
select name, sp.item, price, avg_price
from sp, (select item, avg(price)
          from sp
          group by item) item_avg_price(item, avg_price)
where price <= avg_price and sp.item = item_avg_price.item
order by 1,2,3;
```

name	item	price	avg_price
S1	P1	10	10.0
S1	P2	20	20.0
S2	P3	100	100.0
S3	P4	1000	1000.0
S4	P1	9	10.0

(5 rows)



# Derived Tables - Example 2

- ❖ Find the average supply price for items i.e. avg of avg

```
select avg(avg_price)
from (select item, avg(price)
      from sp
      group by item) item_avg_price(item, avg_price)
;
```

name	item	price	avg_price
S1	P1	10	10.0
S1	P2	20	20.0
S2	P3	100	100.0
S3	P4	1000	1000.0
S4	P1	9	10.0

(5 rows)



# ROLLUP Aggregations

- ❖ Many times we like to “rollup” our aggregates over a hierarchy: e.g. “all items”, “item”, “item, supplier”
- ❖ or country, region, state, city
- ❖ NULLs used in place when group by column is aggregated over: *ALL* in example

```
select item, name, avg(price)
from spo
group by rollup(item, name)
order by item nulls last,
        name nulls last;
```

item	name	avg
P1	S1	10.0
P1	S2	11.0
P1	S4	9.0
P1	<i>ALL</i>	10.0
P2	S1	20.0
P2	<i>ALL</i>	20.0
P3	S2	100.0
P3	<i>ALL</i>	100.0
P4	S3	1000.0
P4	<i>ALL</i>	1000.0
<i>ALL</i>	<i>ALL</i>	191.7

(11 rows)



# NULLs and GROUPING(...)

- ❖ How can one distinguish between NULLs in data and NULLs introduced by ROLLUP?
- ❖ GROUPING(...)
- ❖ also for giving meaning to NULL, 'ALL' works for character columns

```
select item, name, avg(price)
from sp
group by rollup(item, name)
order by item nulls last,
        name nulls last;
```

item	name	avg
P1	S1	10.0
P1	S2	11.0
P1	S4	9.0
P1		
P1		10.0
P2	S1	20.0
P2		20.0
P3	S2	100.0
P3	S4	
P3		100.0
P4	S3	1000.0
P4		1000.0
	S5	
		191.7

(15 rows)



# NULLs and GROUPING(...)

- ❖ How can one distinguish between NULLs in data and NULLs introduced by ROLLUP?
- ❖ GROUPING(...)
- ❖ also for giving meaning to NULL, 'ALL' works for character columns

```
select case when grouping(item) = 1 then 'ALL'
        else item end as item1,
       case when grouping(name) = 1 then 'ALL'
        else name end as name1,
       avg(price)
```

```
from sp
group by rollup(item, name)
order by item, grouping(item),
         name, grouping(name) ;
```

item1	name1	avg
P1	S1	10.0
P1	S2	11.0
P1	S4	9.0
P1		
P1	ALL	10.0
P2	S1	20.0
P2	ALL	20.0
P3	S2	100.0
P3	S4	
P3	ALL	100.0
P4	S3	1000.0
P4	ALL	1000.0
	S5	
	ALL	
ALL	ALL	191.7

(15 rows)



# Extended Grouping

- ❖ ROLLUP most common, but we may not want all levels in report
- ❖ GROUPING SETS lets us specify the exact groups we want
- ❖ On other extreme, CUBE generates every combination of grouping columns ...
- ❖ HAVING works on top

```
select case when grouping(item) = 1 then 'ALL'
        else item end as item1,
       case when grouping(name) = 1 then 'ALL'
        else name end as name1,
       avg(price)
```

```
from sp
group by grouping sets((), (item, name))
order by item, grouping(item),
         name, grouping(name) ;
```

item1	name1	avg
P1	S1	10.0
P1	S2	11.0
P1	S4	9.0
P1		
P2	S1	20.0
P3	S2	100.0
P3	S4	
P4	S3	1000.0
	S5	
ALL	ALL	191.7

(10 rows)



# NULLs Revisited

- ❖ In Tables: unknown or inapplicable values
- ❖ Operations with NULLs result in NULL
- ❖ Introduced by processing
  - ❖ Outer Joins
  - ❖ Extended Grouping

```
select a, b, a+b as aplusb,  
(a > b) as agtb, (a = b) as aeqb  
from t1;
```

a	b	apusb	agtb	aeqb
10	NULL	NULL	NULL	NULL
NULL	NULL	NULL	NULL	NULL
NULL	20	NULL	NULL	NULL
20	30	50	f	f
30	20	50	t	f
(5 rows)				



---

# Operations on NULLs

---

- ❖ Mostly operations with NULLs result in NULL values
- ❖ IS NULL and IS NOT NULL check for NULL values
- ❖ ZEROIFNULL(), COALESCE(), etc convert NULL values to something else.
- ❖ CASE statement is the general case



# NULLs and Aggregates

- ❖ Aggregations ignore nulls for most part, except count(\*)
- ❖ Still all NULL values will result in NULL, except COUNT = 0
- ❖ All NULL group by values are grouped together (as if they are equal)!
- ❖ What if you really wanted to consider two null values equal?

```
select a, sum(b), count(b), count(*)
from t1
group by a;
```

a	sum	count	count
	20	1	2
30	20	1	1
20	30	1	1
10		0	1

(4 rows)



# NULLs

```
select a, b, (a = b) as aeqb,  
        (a = b) OR ((a is null) AND (b is NULL)) as aeqb_incl_null
```

```
from t1;
```

a	b	aeqb	aeqb_incl_null
10			t
	20		
20	30	f	f
30	20	f	f

(5 rows)

- ❖ We can equate nulls if needed
- ❖ Sometimes systems may do that internally when needed



---

# Ordered-Analytics

---

- ❖ Ranking - Find Top 10 rows
- ❖ Time Series
  - ❖ Find moving average of sales
  - ❖ Find cumulative sales for the month for each store
  - ❖ Compare values to a month-ago, year-ago
- ❖ Find ratio to report (monthly sales to annual sales)



# A Simple Rank Query

- ❖ Ranking is one of the common analytic operations.
- ❖ Also used to determine top N rows
- ❖ Uses “olympic” ranking

```
select name, item, price,  
       rank() over (order by price)
```

```
from sp1;
```

name	item	price	rank
S4	P1	9	1
S1	P1	10	2
S2	P1	11	3
S2	P2	20	4
S1	P2	20	4
S2	P3	100	6
S3	P4	1000	7
	P1		8
S5			8
S4	P3		8

(10 rows)



# Ranking “for each” partition

```
select item, name, price,  
       rank() over (partition by item  
                   order by price)
```

```
from sp1;
```

item	name	price	rank
P1	S4	9	1
P1	S1	10	2
P1	S2	11	3
P1			4
P2	S2	20	1
P2	S1	20	1
P3	S2	100	1
P3	S4		2
P4	S3	1000	1
	S5		1

(10 rows)

- ❖ We may want to rank values within partitions (“groups” in aggregation)



# Comparison to Totals

- ❖ Can also report the totals for the “group” or partition across all rows (ratio-to-report, comparison to totals queries)
- ❖ No ORDER BY and No ROWS specification

```
select item, name, price,  
       avg(price) over  
         (partition by item)
```

```
from sp1;
```

item	name	price	avg
P1	S2	11	10.0
P1	S4	9	10.0
P1	S1	10	10.0
P1			10.0
P2	S2	20	20.0
P2	S1	20	20.0
P3	S4		100.0
P3	S2	100	100.0
P4	S3	1000	1000.0
	S5		

(10 rows)



# Moving Sum and Avg

- ❖ Moving Sum / Avg are very common time-series operations

- ❖ Moving Computations

- ❖ ROWS x PRECEDING

- ❖ ROWS BETWEEN  
x PRECEDING and  
Y FOLLOWING

- ❖ Basically any bounded  
“window”

```
select dept, cid, semid, grades,  
       avg(grades) over (order by semid  
                        rows 3 preceding) as mavg
```

```
from e3avg;
```

dept	cid	semid	grades	mavg
CS	564	1	3.000	3.00000
CS	564	2	3.000	3.00000
CS	564	3	4.000	3.33333
CS	564	4	4.000	3.50000
CS	564	5	2.500	3.37500
CS	564	6	3.000	3.37500
CS	564	7	2.667	3.04175
CS	564	8	1.000	2.29175
CS	564	10	3.000	2.41675

(9 rows)



# Cumulative Sum and Avg

- ❖ Cumulative

- ❖ ROWS  
UNBOUNDED  
PRECEDING

- ❖ Basically any  
unbounded  
“window”

```
select dept, cid, semid, grades,  
       avg(grades) over (order by semid  
                        rows unbounded preceding) as cumavg
```

```
from e3avg;
```

dept	cid	semid	grades	cumavg
CS	564	1	3.000	3.000000
CS	564	2	3.000	3.000000
CS	564	3	4.000	3.333333
CS	564	4	4.000	3.500000
CS	564	5	2.500	3.300000
CS	564	6	3.000	3.250000
CS	564	7	2.667	3.166714
CS	564	8	1.000	2.895875
CS	564	10	3.000	2.907444

(9 rows)



---

# Window Functions

---

- ❖ Very verbose
- ❖ `function (<arg>) OVER (  
    [PARTITION BY ...]  
    [ORDER BY ...]  
    <window specification>)`
- ❖ 3 kinds - Ranking, Aggregate, Reference
  - ❖ Ranking - Rank, Row\_number, Percent\_rank
  - ❖ Aggregate - all / most of the aggregate functions
  - ❖ Reference - Lag / Lead (to refer to a particular row in the order)



---

# Window Specification

---

- ❖ MOVING

- ❖ ROWS BETWEEN 3 PRECEDING AND 1 PRECEDING
- ❖ ROWS BETWEEN 3 PRECEDING AND 3 FOLLOWING
- ❖ ROWS 3 PRECEDING = BETWEEN 3 PRECEDING AND CURRENT ROW

- ❖ CUMULATIVE

- ❖ ROWS UNBOUNDED PRECEDING
- ❖ ROWS BETWEEN UNBOUNDED PRECEDING AND 1 PRECEDING

- ❖ “TOTAL”

- ❖ none = ROWS BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING



# What about HAVING?

- ❖ Unfortunately, the SQL standard didn't provide for HAVING equivalent

```
select item, name, price, myrank
from (select item, name, price,
            rank() over (partition by item
                        order by price)
            as myrank
```

```
      from sp1) as d1
where myrank = 1;
```

item	name	price	myrank
P1	S4	9	1
P2	S2	20	1
P2	S1	20	1
P3	S2	100	1
P4	S3	1000	1
	S5		1

(6 rows)



# Nesting Aggregates

- ❖ Window Functions can have nested aggregates inside
- ❖ Equivalent to the semantics on the right

```
select item,  
       cast(avg(price) as decimal(5,1)),  
       rank() over (order by avg(price))  
         as myrank
```

```
from sp1  
group by item;
```

item	avg	myrank
P1	10.0	1
P2	20.0	2
P3	100.0	3
P4	1000.0	4
		5

```
select item, avg_price,  
       rank() over (order by avg_price)  
         as myrank
```

```
from (select item,  
            avg(price) as avg_price  
      from sp1  
      group by item) item_avg_price;  
item | avg_price | myrank
```

P1	10.000000000000000000	1
P2	20.000000000000000000	2
P3	100.000000000000000000	3
P4	1000.000000000000000000	4
		5

(5 rows)