

SQL in Programs

Basic Architecture

- ❖ Application Layer - what most users see, talks SQL
- ❖ Parsing / Planning Layers - the intelligence
- ❖ Runtime or execution Layer - the brawn
- ❖ Storage Layer - where data resides, may include simple access layer

Applications

Parsing


Planning

Processing

Data Access

Data in SSD / HDD

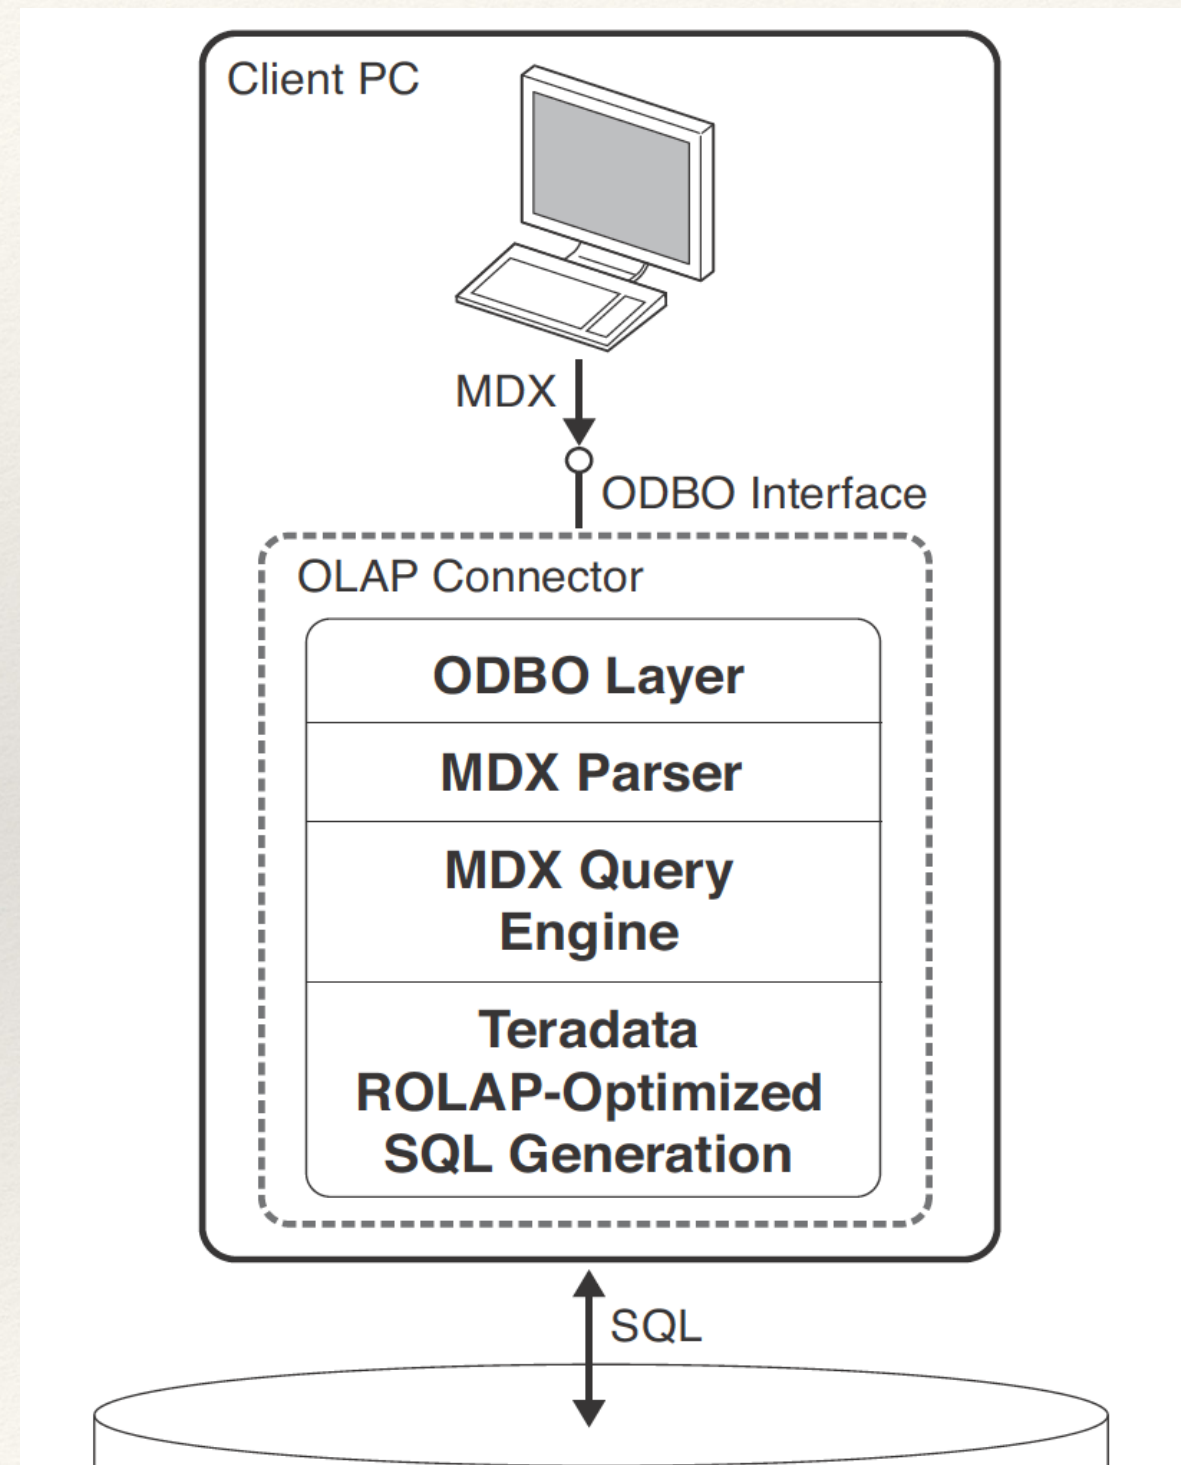
Typical Connections to SQL



- ❖ Applications connect via Embedded SQL, ODBC, JDBC, SQLJ, CLI
- ❖ Middle Tier: e.g. MicroStrategy, BusinessObjects, etc.
- ❖ Client SQL Utilities: psql, Teradata® Bteq, Teradata® SQL Assistant

Connecting Excel to Teradata®

- ❖ Excel “talks” MDX/ODBO
- ❖ The Connector takes MDX and converts it into SQL (in the client)
- ❖ It connects to Teradata® using ODBC
- ❖ SQL gets executed by Teradata®



Basic Framework

- ❖ SQL calls accessible from a host language, e.g. C or Java
- ❖ From the host language one can:
 - ❖ Connect to data-source
 - ❖ Execute SQL
 - ❖ Fetch results
 - ❖ Process results

Different API's/Standards

- ❖ Embedded SQL standard
- ❖ CLI - or Call Level Interface - basis for ODBC
 - ❖ Teradata® Bteq (similar to Postgres psql)
- ❖ ODBC - Open Database Connectivity
 - ❖ Teradata® SQL Assistant, MS Sql Server Mgmt Studio
- ❖ JDBC - Java (Open) Database Connectivity
- ❖ SQL/J - Sort of Combining Java with Embedded SQL

Overview

- ❖ May need to describe the environment the system is running in, e.g. ODBC version
- ❖ Connect (login) to the DBMS, and then
 - ❖ Execute the SQL's we want (usually need to build it first)
 - ❖ Fetch the results
 - ❖ Post-process and repeat
- ❖ Disconnect from the DBMS

Embedded SQL

- ❖ SQL calls are embedded in a host-language, e.g. C
- ❖ It's standardized, so same source code will (mostly) work with typical DBMS
- ❖ SQL interface precedes with EXEC SQL
- ❖ A preprocessor converts the EXEC SQL to calls to the database system
- ❖ The final program is compiled as executed with integrated SQL processing

Environment & Connecting

```
EXEC SQL CONNECT TO db1@localhost;
```

- ❖ The source code is “preprocessed” by a DBMS specific preprocessor, e.g. `ecpg` for postgres.
- ❖ depending on the complexity of connection (local, remote) and security setup, more information may be needed by the `CONNECT TO`.

Shared Variables

```
EXEC SQL BEGIN DECLARE SECTION;  
    int sid;  
    varchar name[21];  
EXEC SQL END DECLARE SECTION;
```

- ❖ To communicate, the host language and the SQL share variables
- ❖ The data types much match between the two. This is usually well documented and is fairly intuitive, but can vary across systems

Accessing Results via Cursor

```
EXEC SQL DECLARE foo CURSOR FOR SELECT sid, last_name FROM s2;  
EXEC SQL OPEN foo;  
for (i = 0; i < 10; i++) {  
    EXEC SQL FETCH NEXT FROM foo INTO :sid, :name;  
    printf("%d\t%s\n", sid, name.arr);  
}  
EXEC SQL CLOSE foo;
```

- ❖ Cursor gives a row-by-row access to the query results
- ❖ Then one can fill in the host variables with the row values
- ❖ Finally, one should close the cursor
- ❖ Cursors can be “updatable” when accessing base tables or updatable views

Put It Together

```
#include <stdio.h>
int main()
{
    int i;
    EXEC SQL BEGIN DECLARE SECTION;
        int sid;
        varchar name[21];
    EXEC SQL END DECLARE SECTION;
    EXEC SQL CONNECT TO db1@localhost;
    EXEC SQL DECLARE foo CURSOR FOR
        SELECT sid,last_name
        FROM s2 ORDER BY sid;
    EXEC SQL OPEN foo;
    for (i = 0; i < 10; i++) {
        EXEC SQL FETCH NEXT FROM foo
            INTO :sid, :name;
        printf("%d\t%s\n", sid, name.arr);
    }
    EXEC SQL CLOSE foo;
    EXEC SQL DISCONNECT all;
}
```

```
$ ecpg pg1.pgc
$ gcc -o esql pg1.c -lecpg
$ esql
1      Washington
16     Lincoln
32     Roosevelt
33     Truman
34     Dwight
35     Kennedy
36     Johnson
37     Nixon
38     Ford
39     Carter
```

ODBC/JDBC

- ❖ A “driver” model introduces another layer of indirection
- ❖ The “client” program talks to the driver that then connects to the data-source
- ❖ Client program can be fairly independent of the data-source
- ❖ Any data-source having an ODBC/JDBC driver interface can be used (and needn't be an RDBMS)

JDBC

- ❖ Java Database Connectivity
- ❖ Like ODBC but for Java
- ❖ Drivers are available for most databases, e.g. Postgres
 - ❖ Provided by DBMS or 3rd party
- ❖ JDBC-to-ODBC bridge makes ODBC data-sources accessible

Environment and Connection

```
Class.forName("org.postgresql.Driver"); //environment
Connection c = DriverManager.getConnection
    ("jdbc:postgresql://localhost:5432/db1",
     "ambuj", "123");
```

- ❖ First thing is to load the needed Driver
- ❖ Then get the connection to “db1” database on postgres
- ❖ Logic “ambuj”, password is irrelevant here (depends on security)
- ❖ General case below

```
Connection c =
DriverManager.getConnection("jdbc:somejdbcvendor:otherdataneededbyvendor",
    "myLogin", "myPassword");
```

Create A SQL Statement

```
stmt = c.createStatement();  
String sql = "CREATE TABLE MyNewTbl " +  
             "(ID INT PRIMARY KEY     NOT NULL," +  
             " NAME                   TEXT     NOT NULL," +  
             " AGE                    INT      NOT NULL," +  
             " ADDRESS                CHAR(50), " +  
             " SALARY                 REAL)";  
stmt.executeUpdate(sql);  
stmt.close();
```

- ❖ We must create a statement
- ❖ And create a SQL that would be part of the statement
 - ❖ CREATE TABLE in our example
- ❖ executeUpdate - for DDL as no data returned

Insert A Row

```
c.setAutoCommit(false);
stmt = c.createStatement();
String sql = "INSERT INTO MyNewTbl(ID,NAME,AGE,ADDRESS,SALARY) "
            + "VALUES (1, 'Paul', 32, 'California', 20000.00 );";
stmt.executeUpdate(sql);
sql = "INSERT INTO MyNewTbl(ID,NAME,AGE,ADDRESS,SALARY) "
      + "VALUES (2, 'Allen', 25, 'Texas', 15000.00 );";
stmt.executeUpdate(sql);
...
stmt.close();
c.commit();
```

- ❖ To Insert a Row, same process but difference SQL
- ❖ We can insert many individual rows, using different SQL but same statement
- ❖ Commit after all the inserts (otherwise it does after each one)

What about SELECT

```
stmt = c.createStatement();
ResultSet rs = stmt.executeQuery( "SELECT * FROM MyNewTbl;" );
while ( rs.next() ) {
    int id = rs.getInt("id");
    String name = rs.getString("name");
    int age = rs.getInt("age");
    String address = rs.getString("address");
    float salary = rs.getFloat("salary");
    System.out.println( "ID = " + id );
    System.out.println( "NAME = " + name );
    System.out.println( "AGE = " + age );
    System.out.println( "ADDRESS = " + address );
    System.out.println( "SALARY = " + salary );
    System.out.println();
}
rs.close();
stmt.close();
```

- ❖ What if we don't know all the data types?


```
$ java PGSelect  
Opened database successfully  
ID = 1  
NAME = Paul  
AGE = 32  
ADDRESS = California  
SALARY = 20000.0
```

```
ID = 2  
NAME = Allen  
AGE = 25  
ADDRESS = Texas  
SALARY = 15000.0
```

```
ID = 3  
NAME = Teddy  
AGE = 23  
ADDRESS = Norway  
SALARY = 20000.0
```

```
ID = 4  
NAME = Mark  
AGE = 25  
ADDRESS = Rich-Mond  
SALARY = 65000.0
```

```
Operation done successfully
```


Metadata About A Query

```
ResultSet rs = stmt.executeQuery( "SELECT * FROM MyNewTbl;" );
ResultSetMetaData rsmd = rs.getMetaData();
for (int i = 1; i < rsmd.getColumnCount(); i++){
    System.out.println(rsmd.getColumnTypeName(i)+"\t\t"
        +rsmd.getColumnName(i));
}
```

- ❖ Can get the output
Column Types and
Names

```
$ java PGMeta
Opened database successfully
int4    id
text    name
int4    age
bpchar  address
float4   salary
Operation done successfully
```

Motivations for SQL Programming

- ❖ Interacting with applications that work with data
- ❖ Presenting a different view of the data in the DBMS
 - ❖ Many BI apps fall in this category
- ❖ Doing things that are hard to do in SQL
 - ❖ Very common in old days but less so

Issues

- ❖ Performance - data fetched serially
- ❖ Scale - result dataset too large
- ❖ If we want to further use the DBMS, putting data back is equally slow

Solutions

- ❖ Move computation to the data
- ❖ New DBMS features
 - ❖ Computations built into the system
 - ❖ Top N eliminates need for ORDER BY and fetch
- ❖ Temp Tables, compute and store results in the DBMS
 - ❖ Multi-step processing can be done in the DBMS
- ❖ Extensibility features, move computations to the DBMS
 - ❖ scoring, statistical analysis