

B⁺-Tree

Needed Capabilities

- ❖ quickly scan all data
- ❖ find exact (equality) matched data
- ❖ do efficient range searches
- ❖ do efficient joins
- ❖ do efficient insert / delete / updates
- ❖ For SSD / HDD resident data \Rightarrow fewest I/O's

Fast Lookup

- ❖ Hash Tables
- ❖ Balanced Binary Trees
- ❖ B+ Tree is balanced binary tree for disks

What are we storing?

- ❖ In leaf nodes we can store
 - ❖ rows (record) for the primary index
 - ❖ need not be same as primary key
 - ❖ index_value, *rid* for a unique secondary index
 - ❖ e.g. for candidate keys, or even primary key
 - ❖ index_value, *rid*-list for secondary index

Primary Index

- ❖ Clustered, Sorted
- ❖ Choice of Primary Index should be done carefully
 - ❖ study the workload
 - ❖ joins and range searches are specially valuable
 - ❖ primary key can make a good primary index and should be first (default) choice

B+ tree nodes

- ❖ Optimized for disk-resident data
- ❖ Keep pages reasonably full
- ❖ B+ tree nodes are disk pages
- ❖ Two kinds
 - ❖ Internal
 - ❖ Leaf


Internal Nodes

- ❖ Keep the index_values (keys) and “pointers” to pages
- ❖ $(p_0, key_1, p_1, \dots, key_n, p_n)$
- ❖ p_i is the pointer to subtree containing keys between key_i and key_{i+1} ; usually $> key_i$ and $\leq key_{i+1}$
- ❖ p_0 points to the subtree all keys $\leq key_1$
- ❖ p_n points to the subtree all keys $> key_n$

Leaf Node

- ❖ contains keys with data, such as
 - ❖ integrated record / rows i.e. keys are part of the record
 - ❖ set of (key, record) pairs
 - ❖ set of (key: list of *rids*)
 - ❖ set of (key, *rid*) pairs

B+ tree



- ❖ $\text{index_val} = \text{key}$
- ❖ 2 indicates record with key value of 2.


Properties

- ❖ Every node (except possibly root) is at least $1/2$ full
 - ❖ if max entries is 2^*d , then must have at least d entries
 - ❖ d is called order of the tree
- ❖ Every leaf is at the same distance from root
 - ❖ tree is balanced to avoid bad cases

Properties (contd.)

- ❖ Search at $\log_F(N)$ cost
 - ❖ F is the fanout (at least d)
 - ❖ N is the number of leaf pages
- ❖ Insert / delete at least $\log_F(N) + 1$, possibly a little more
 - ❖ 1 more I/O to write the updated leaf page
 - ❖ may need additional I/O to update interior nodes

Simple Example



- ❖ search for PI value = 7
- ❖ search for PI value = 18
- ❖ search for PI value > 31

B+ trees



- ❖ For a 4K page, 8 byte keys, 8 byte “pointer”
- ❖ Typical $\text{Order}(d) = 100$, i.e. $\max(2d) = 200$
- ❖ average fanout = 133
- ❖ So at height of 4 we have $133^4 = 313\text{M}$ records
- ❖ height of 3: $133^3 = 2.3\text{M}$ records
- ❖ Often hold top levels in buffer pool
- ❖ Larger page sizes \Rightarrow more fanout

Inserting Data

- ❖ Find correct leaf L
- ❖ Put data entry onto L
 - ❖ If L has enough space, done!
 - ❖ Else, must **split** L (into L and a new node L2)
 - ❖ Redistribute entries evenly, **copy up** middle key.
 - ❖ Insert internal (index) entry pointing to L2 into parent of L.
- ❖ This can happen recursively
 - ❖ **To split (internal) index node**, redistribute entries evenly, but **push up** middle key. (Contrast with leaf splits.)
- ❖ Splits “grow” tree horizontally; root split increases height.
 - ❖ Tree growth: gets wider or one level taller at top.


Insert into Leaf

- ❖ If there is space in page, records are moved to create a space in the proper order
- ❖ The record is then inserted in page




Insert into Leaf

- ❖ Observe how minimum occupancy is guaranteed in both leaf and index page splits.
- ❖ The “middle valued” key is **copied-up**.
- ❖ It happens to be the value of inserted key (with record) in this case




Node splits and 6 is copied up



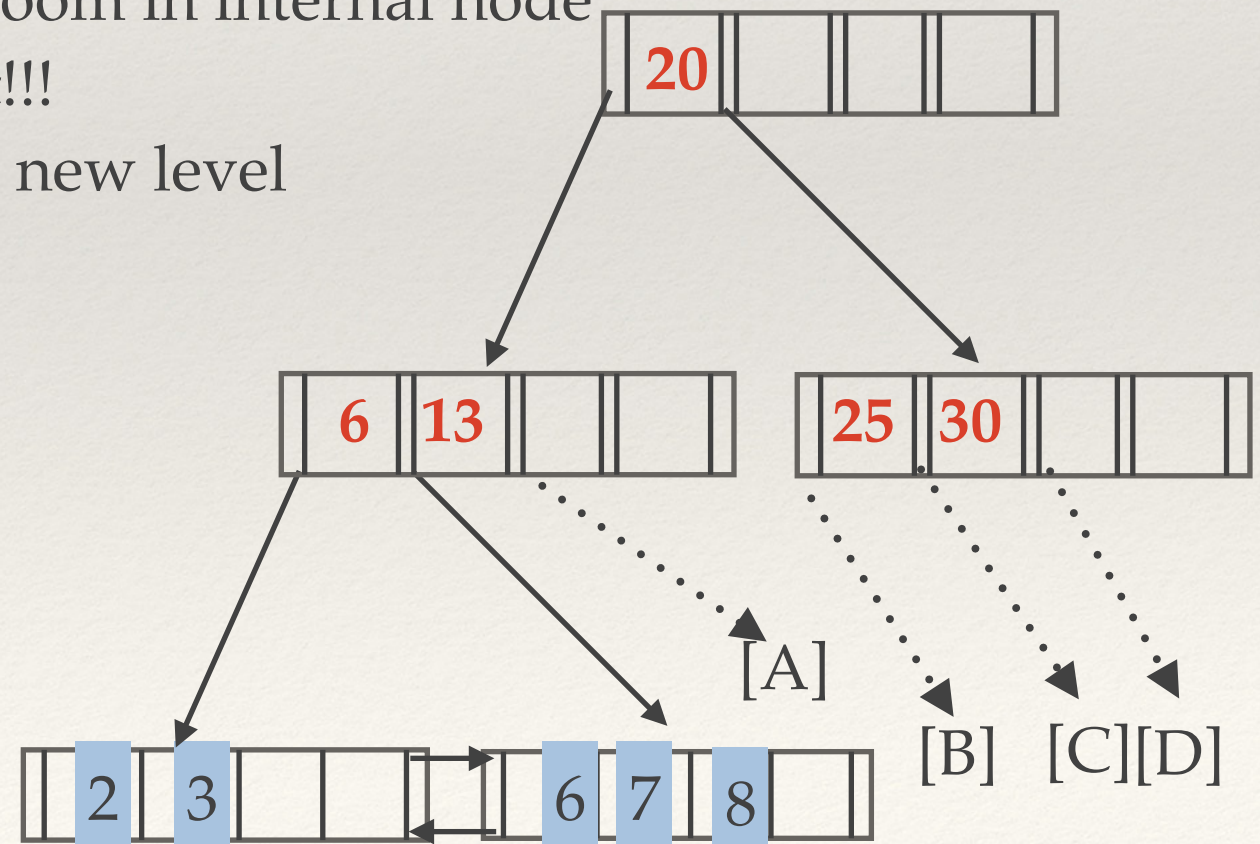
Insert into Internal Node

- ❖ Observe how minimum occupancy is guaranteed in both leaf and index page splits.
- ❖ If there is a split, the “middle valued” key is **pushed-up**
- ❖ New level may be needed

❖ Insert record with key 6



- ❖ Node splits and 6 is copied up
- ❖ No room in internal node
- ❖ Split!!!
- ❖ Add new level




Deleting an entry from B+ Tree

- ❖ Start at root, find leaf L where entry belongs.
- ❖ Remove the entry.
 - ❖ If L is at least half-full, done!
 - ❖ If L has only $d-1$ entries,
 - ❖ Try to re-distribute, borrowing from sibling (adjacent node with same parent as L).
 - ❖ If re-distribution fails, merge L and sibling.
- ❖ If merge occurred, must delete entry (pointing to L or sibling) from parent of L.
- ❖ Merge could propagate to root, decreasing height.


Delete record with Key 6

❖ Note that delete is not symmetric with insert



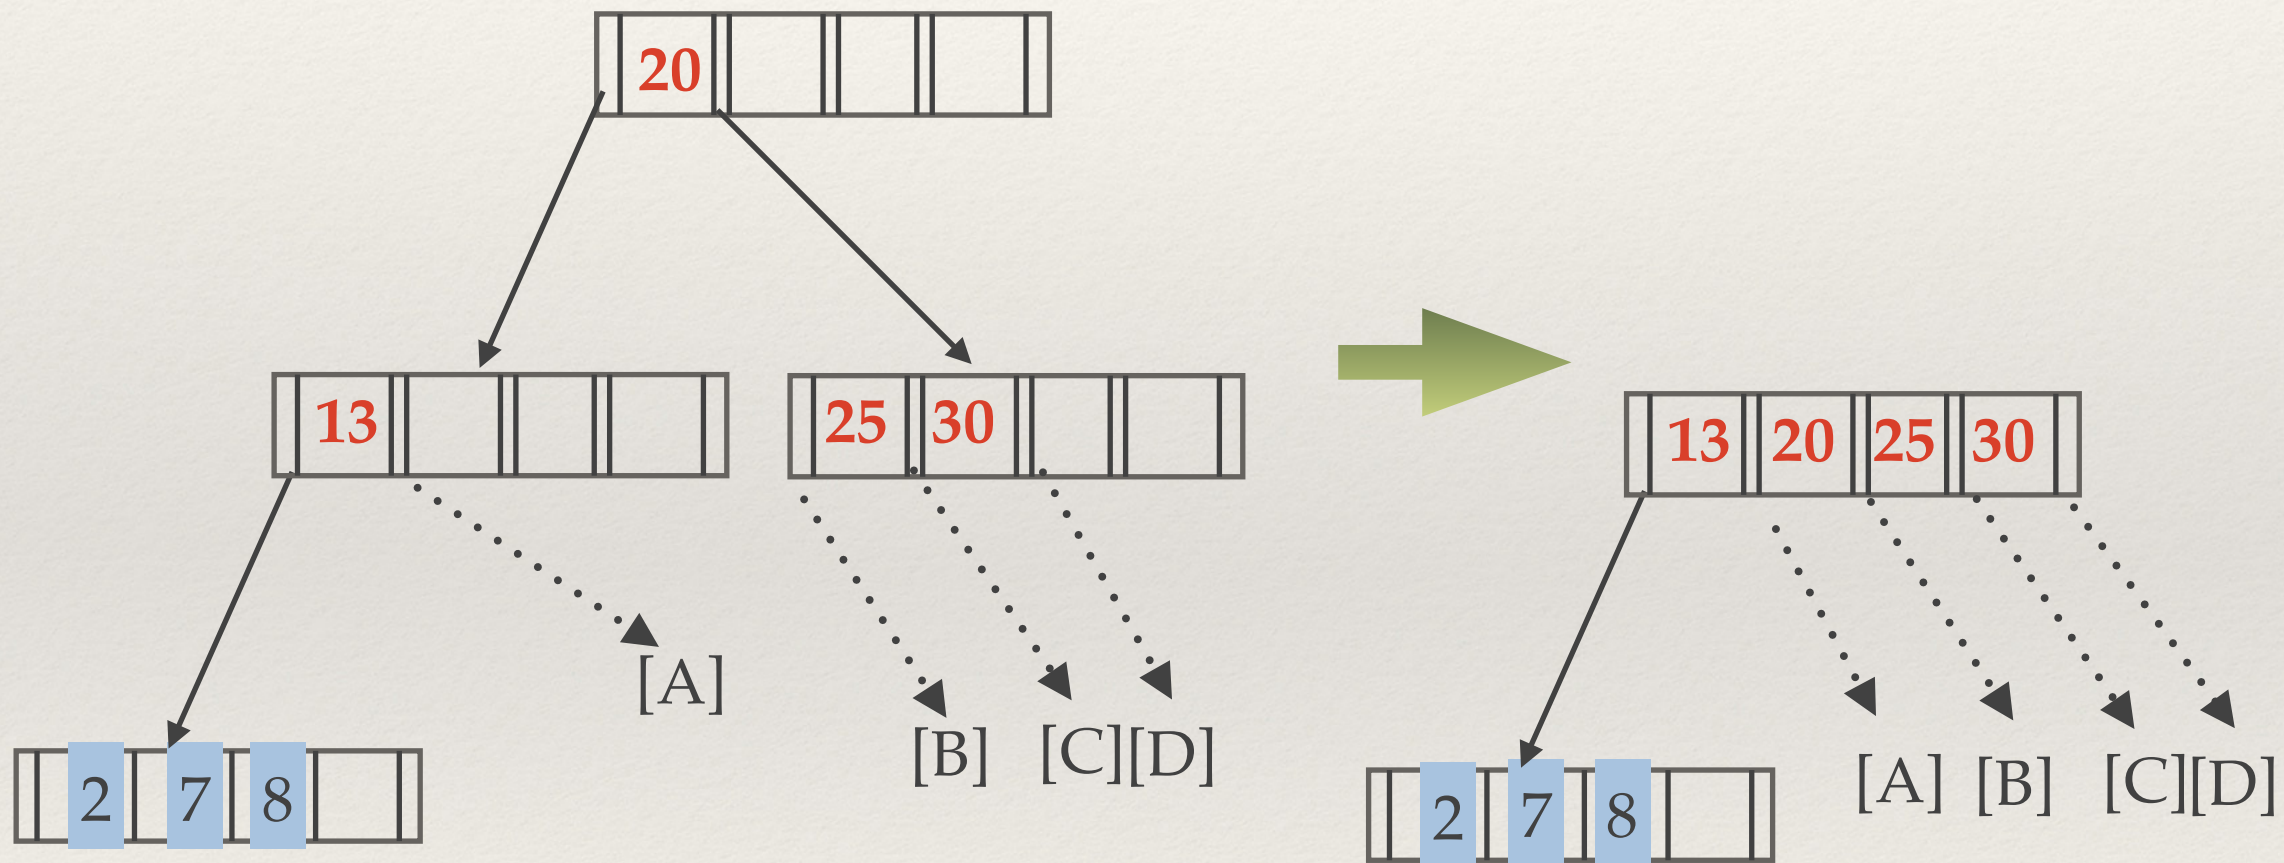
Delete record with Key 3

- ❖ No redistribution possible (both only 1 / 2 full)
- ❖ merge with sibling
- ❖ But what about the internal node with 13?



Delete record with Key 3

- ❖ No redistribution possible (both only 1 / 2 full)
- ❖ merge with sibling
- ❖ If propagated to root, remove a level




Deletion contd.

- ❖ What about a different order?

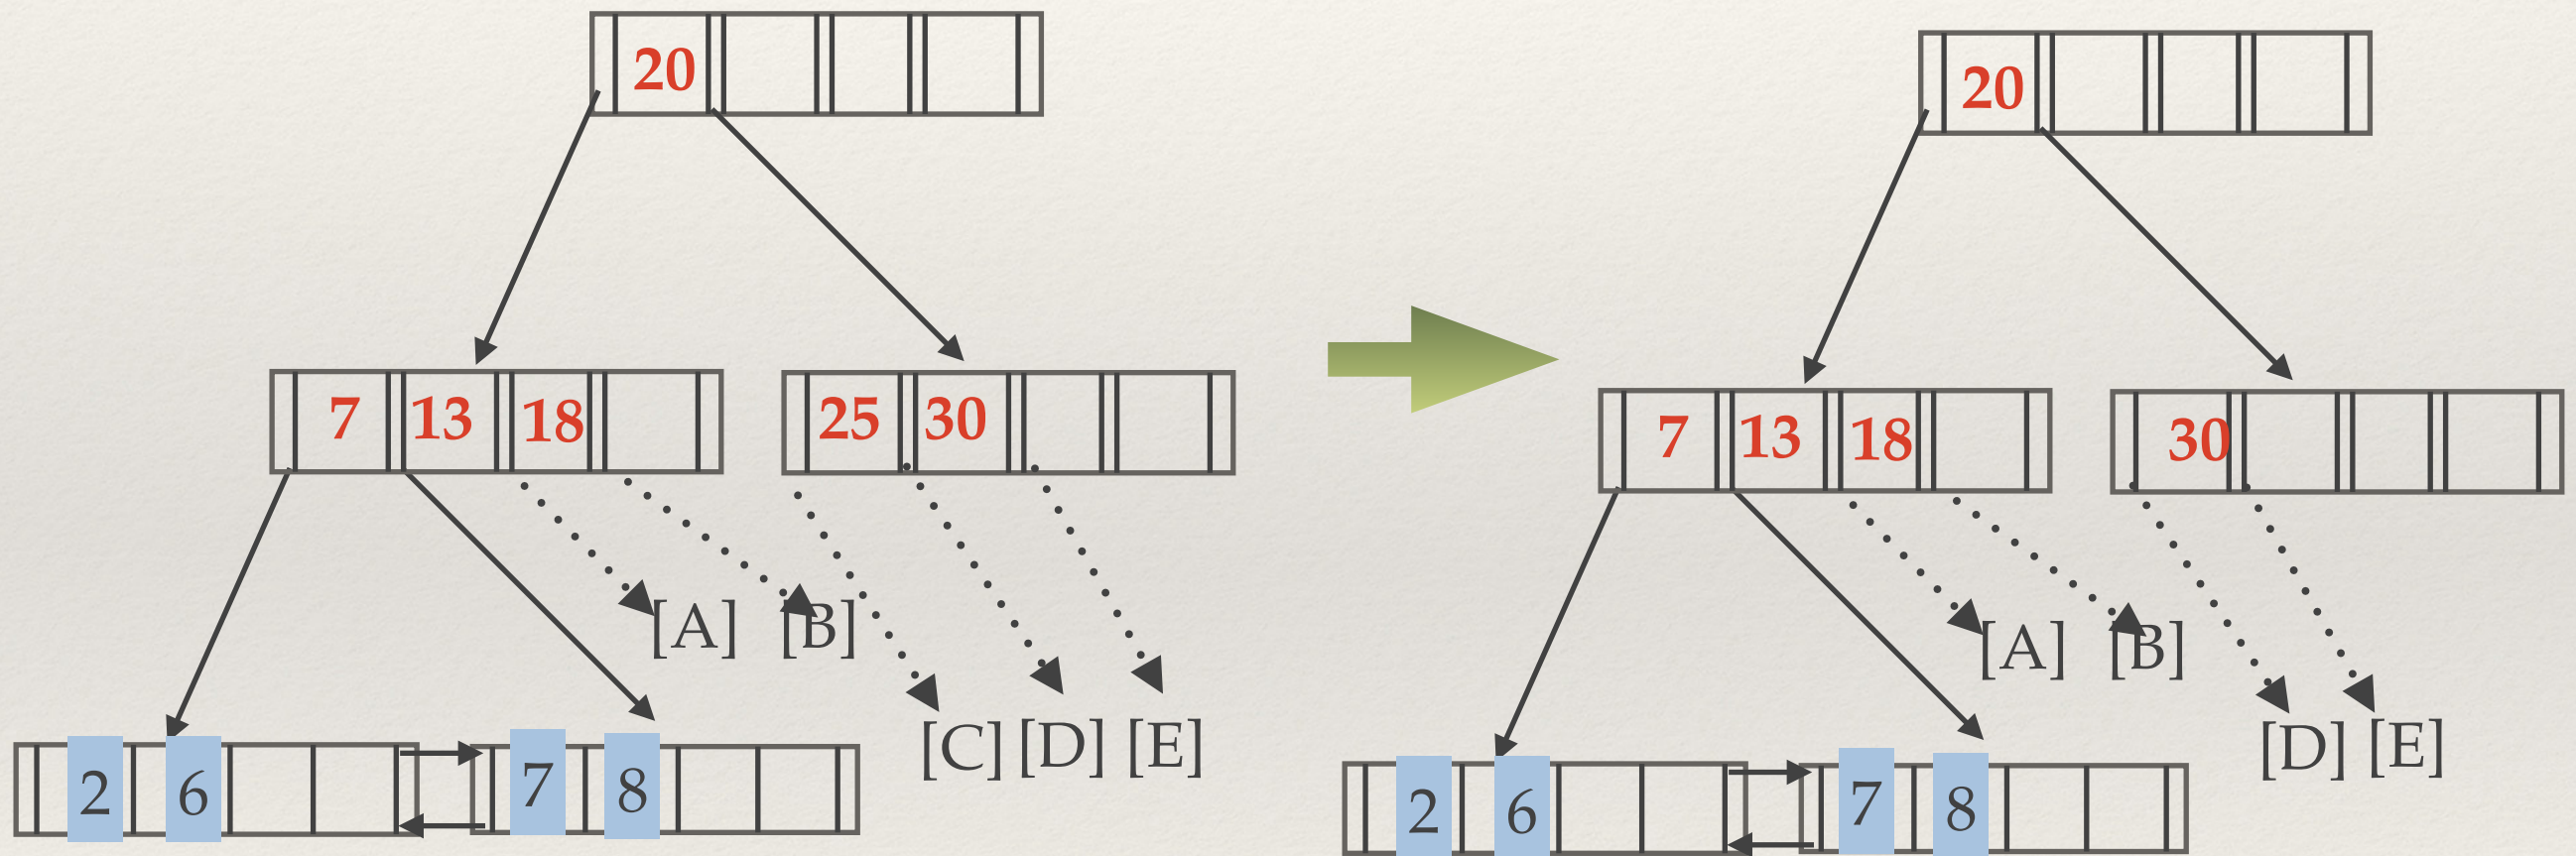
Delete record with Key 3

- ❖ Redistribute from sibling
- ❖ Adjust the key (copy-up from sibling here)




Delete all records from page [C]

- ❖ the right internal node is empty
- ❖ delete the key from internal
- ❖ borrow from sibling




Borrowing from Sibling - Internal Node

- ❖ Entries are redistributed by **pushing through** the splitting entry in the parent node




Prefix Key Compression

- ❖ Since the internal (index) nodes only “direct” search can just use prefix for longer strings, say first three letters of state names: MAS, MIC, MIS, MIS, MON, etc.
- ❖ Need to ensure that index entry is $>$ all (uncompressed) key values to left
- ❖ Insert/Delete need additional care




Bulk Loading into a B+ Tree

- ❖ If we have a large collection of records, and we want to create a B+ tree on some field, doing so by repeatedly inserting records is very slow.
- ❖ Bulk Loading can be done much more efficiently.
- ❖ Initialization: Sort all data entries, insert pointer to first (leaf) page in a new (root) page.




Bulk Loading into a B+ Tree

- ❖ Index entries for leaf pages always entered into right-most index page just above leaf level. When this fills up, it splits. (Split may go up right-most path to the root.)
- ❖ Much faster than repeated inserts, especially when one considers locking!



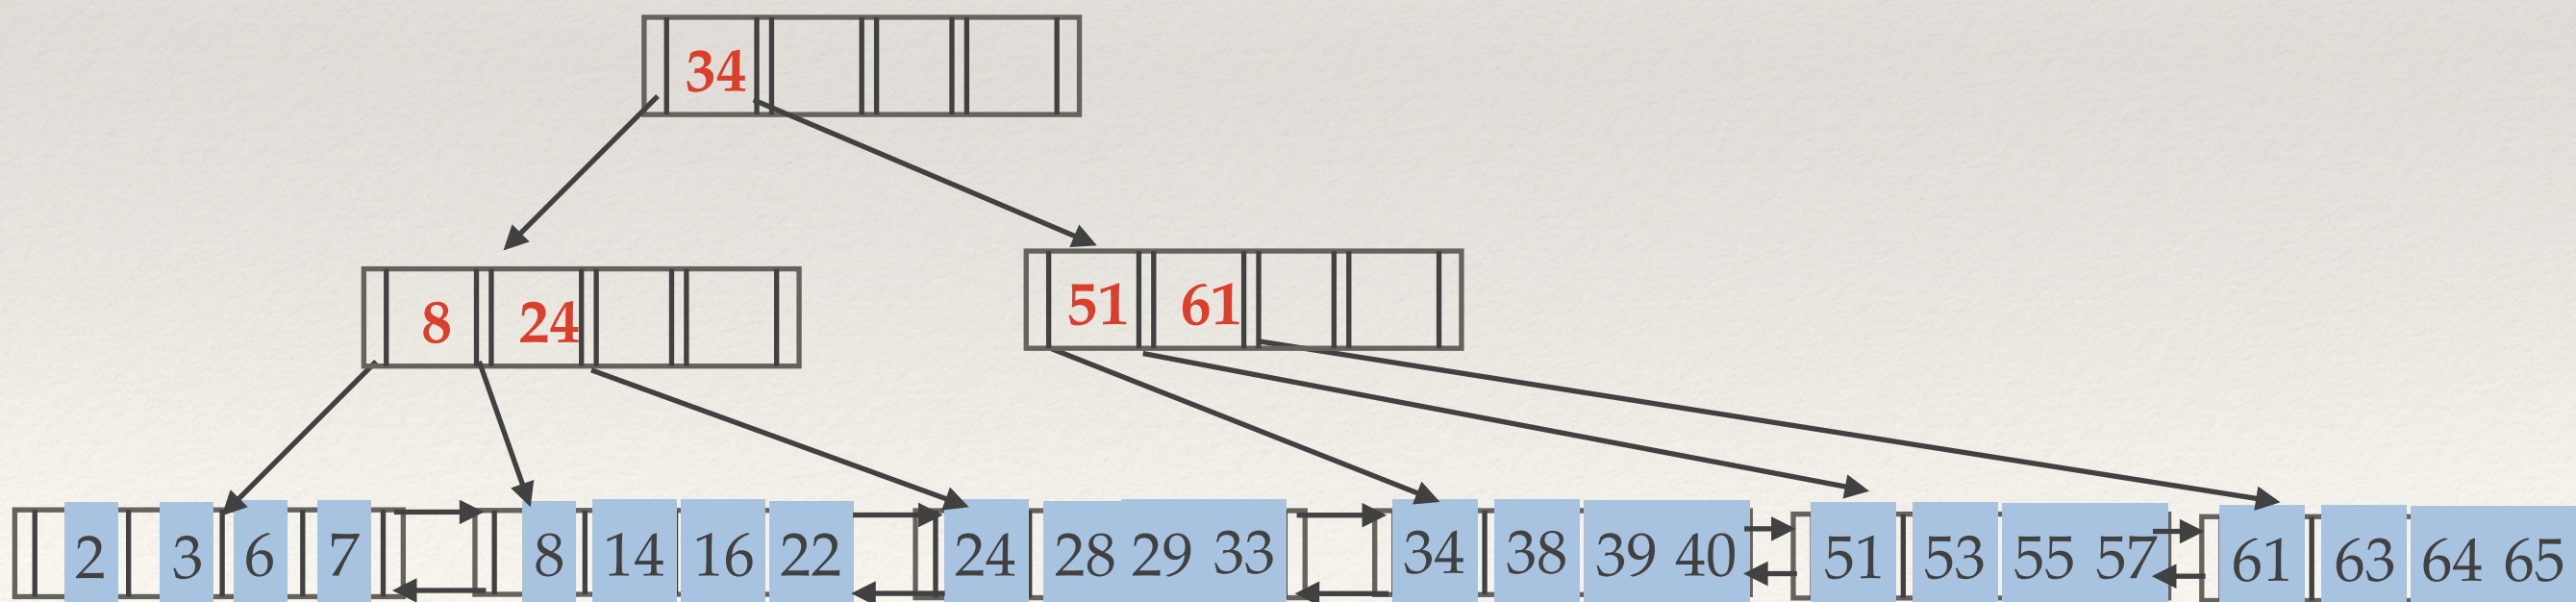
Bulk Loading into a B+ Tree

- ❖ Index entries for leaf pages always entered into right-most index page just above leaf level. When this fills up, it splits. (Split may go up right-most path to the root.)
- ❖ Makes for a compact (fuller) tree



Bulk Loading into a B+ Tree

- ❖ Index entries for leaf pages always entered into right-most index page just above leaf level. When this fills up, it splits. (Split may go up right-most path to the root.)
- ❖ Makes for a compact (fuller) tree



Bulk Loading

- ❖ Option 1: multiple inserts.
 - ❖ Slow
 - ❖ Does not guarantee sequential storage of leaves
- ❖ Option 2: Bulk Loading
 - ❖ Has advantages for concurrency control.
 - ❖ Fewer I/Os during build.
 - ❖ Leaves likely to be stored sequentially
 - ❖ Can control “fill factor” on pages

How many pointers?

- ❖ “order” or minimum number of pointers on node is usually replaced with “page half full” criterion
- ❖ leaves may contain variable number of entries due to variable size records

Review

- ❖ B+ Trees are ideal data structure for most work loads
- ❖ Bulk loading makes it faster and that's what is done as much as possible
- ❖ Widely used and optimized, including Teradata®
- ❖ Heap Files may be good if no equality or range search