

# Indexing Techniques



---

# Storing Tables

---

- ❖ B+ Trees store records with keys
- ❖ Keys should not have too many duplicates i.e. as unique as possible, though uniqueness is not required
- ❖ Normally, in leaf node, the key is just part of record and no duplicate copy is stored



---

# Table to B+Tree

---

```
create table courses(  
  dept varchar(10),  
  cid smallint,  
  name varchar(20),  
  credits smallint,  
  description varchar(1000),  
  last_taught varchar(20),  
  primary key(dept, cid));
```

- ❖ This would be stored as a B+ tree with the *key* (dept, cid)
- ❖ Index (interior) nodes will have key values and pointers
- ❖ Leaf nodes will just have the records ordered by key.
- ❖ DBMS can easily extract keys (or any attributes) from records



---

# Secondary Index as a Table

---

- ❖ We can think of secondary index as a table of type
  - ❖  $\text{Table\_SI}(\text{index attribute}(s), \text{rid-list})$  or
  - ❖  $\text{Table\_USI}(\text{index attribute}(s), \text{rid})$
- ❖ The structure of this table would be exactly as expected
  - ❖ B+ tree with index\_attributes in index nodes and the index\_attributes + rid(s) in the leaf nodes



---

# Hash-Based Indexing

---

- ❖ Tree and Hash based indexing are two main approaches
- ❖ B+ Tree is the main kind for tree-based indexing
- ❖ On to Hash-based indexing ...



---


# Static Hash Based Storage

---

- ❖ Let  $M$  be the expected number of pages in the file (table)
- ❖  $\text{hash}(\text{key\_val}) \bmod M = \text{bucket}$  to which record with  $\text{key} = \text{key\_val}$  would be allocated
- ❖ If the initially allocated page for the bucket is full, then add more pages to that bucket
  - ❖ like a heap file, but per bucket



# Static Hashing





---

# Static Hashing

---

- ❖ Hash function chosen for even distribution
- ❖ Long chains can form quickly when storing regular data
- ❖ Not adjustable to data volumes
- ❖ BUT
  - ❖ for non-unique secondary indices i.e. (key, *rid-list*) kind, the number of buckets is sort-of fixed, so it would work



---


# Extendible Hashing

---

- ❖ When bucket (the primary page) becomes full, why not split the page (like B+ tree)?
- ❖ We can double the number of buckets...
  - ❖ by introducing another level of indirection
- ❖ Instead of buckets, the hash function points to a directory (of pointers entry).




# Static Hashing with Directory



$$\text{hash}(key) = f(key) \text{ MOD } 4$$




# Static Hashing with Directory



$\text{hash}(key) = f(key) \text{ MOD } 2^{32}$  but using last 2 bits



# Expanding Directory

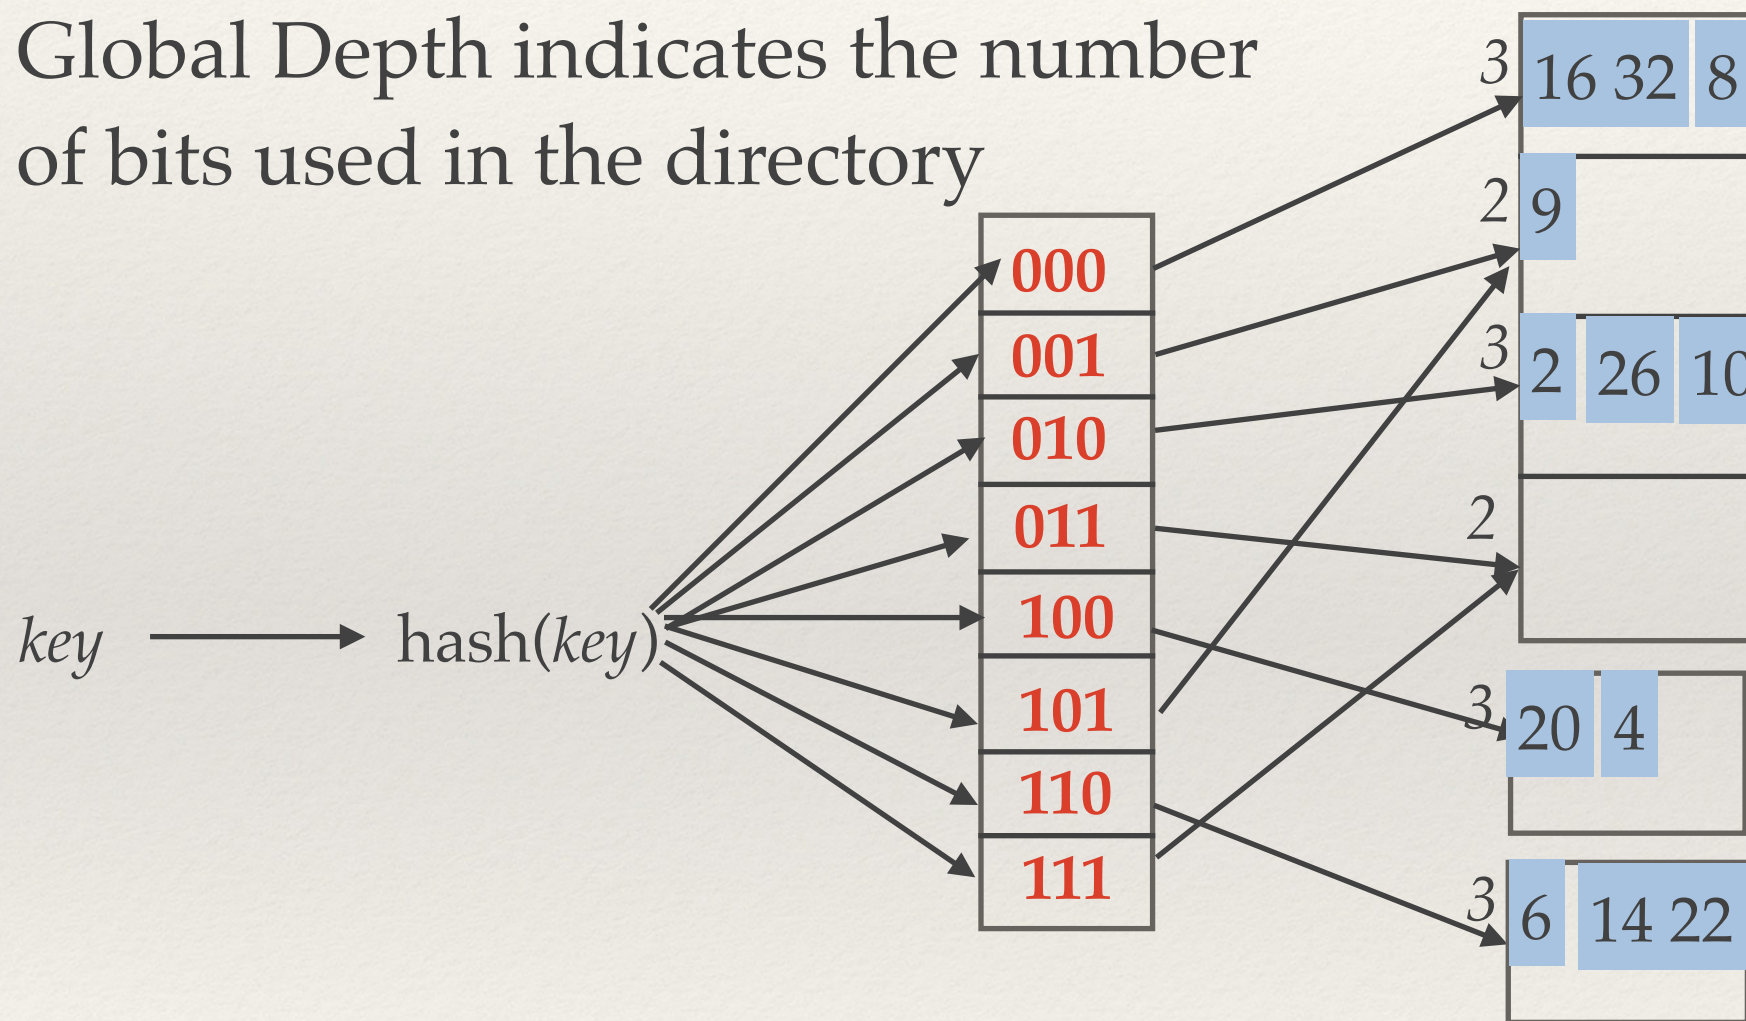


$\text{hash}(\text{key}) = f(\text{key}) \text{ MOD } 2^{32}$  but now using last 3 bits



# Extendible Hashing

- ❖ *LOCAL DEPTH* indicates the number of bits used by the bucket
- ❖ Global Depth indicates the number of bits used in the directory



$\text{hash}(key) = f(key) \text{ MOD } 2^{32}$  but now using last 3 bits



---

# Local and Global Depth

---

- ❖ Last  $n$  bits of hash tell us that the *key* belongs to which directory
  - ❖ For directory: use #bits = global depth
  - ❖ For bucket: use #bits = local depth (entry  $\in$  bucket)



---


# Bucket Splits

---

- ❖ Bucket may need to split upon insert
- ❖ If local depth < global depth, add a new bucket and move records to new bucket per the extra bit of hash (local depth now increases by 1)
- ❖ If they are same then the directory needs to double (simply copy it and then “fix” the pointers to new page)



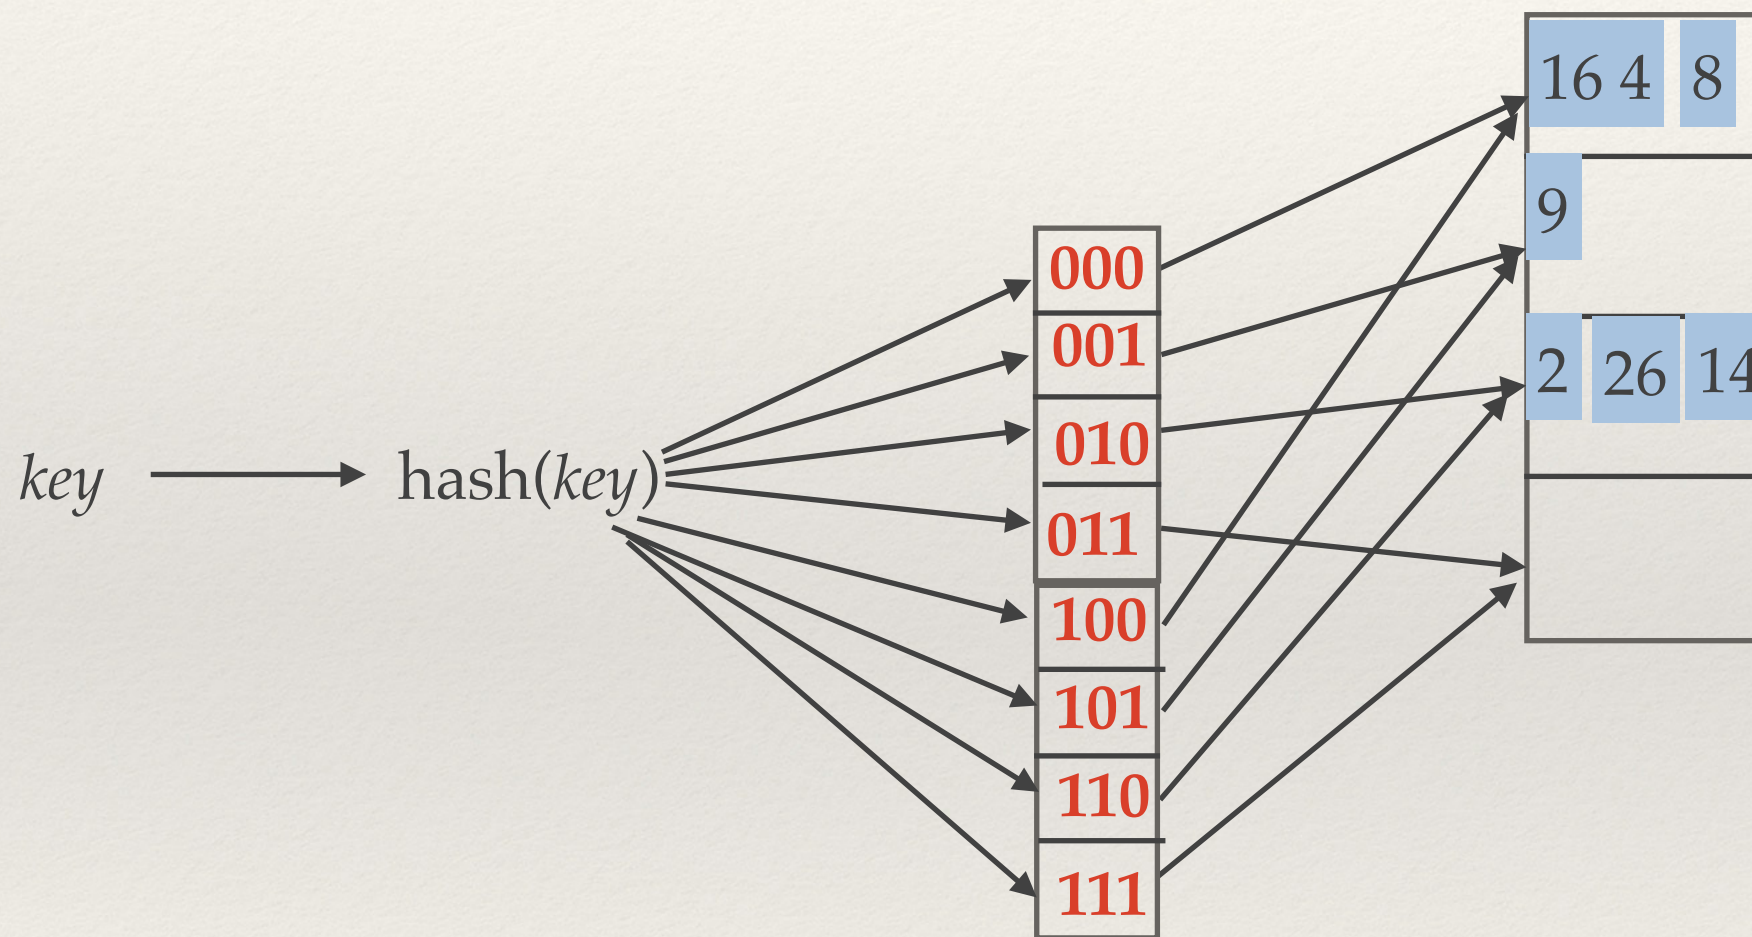
# Example



- ❖ Record **10** needs to be inserted in bucket “**10**”
- ❖ Local depth = global depth = 2




# Example: Doubling of directory



- ❖ Directory split by copying
- ❖ global depth is now 3



# Example: Moving Data and Adjustment



- ❖ Bucket 10 is split into 010 and 110
- ❖ Records are redistributed
- ❖ Local depth is increased by 1 for split buckets
- ❖ Pointer to new bucket is fixed



---

# Extendible Hashing

---

- ❖ If directory is in memory, equality search still takes only one I/O. (Otherwise, 2)
- ❖ Upon delete if the bucket is empty, it can be merged with its “partner” (ignoring the first bit of hash)
- ❖ If all directory elements point to the same bucket as its split partner, can halve the directory too



# Bitmap Index

rid	dept	cid	sid	semid	instructor	grade
1	CS	564	2016001	10	Codd	A
2	CS	564	2012144	10	Codd	C
3	CS	564	2014101	10	Codd	A
4	CS	367	2014101	8	Aho	A
5	CS	367	2012144	8	Aho	C
6	CS	367	2015001	8	Aho	B
7	PolSci	104	2015001	10	Madison	A
8	PolSci	104	2014101	10	Madison	A
9	PolSci	104	2012144	10	Madison	B
10	PolSci	104	2012144	8	Franklin	F
11	Math	234	2012144	1	Newton	D
12	Math	201	2012144	2	Laplace	D

- ❖ Imagine a making an index on dept
- ❖ Represent the rid's as a bitmap
- ❖ 8 byte *rid*, 12 rows => 64 vs. 6 bytes

dept	rid-list
PolSci	7,8,9,10
CS	1,2,3,4,5,6
Math	11,12

dept	rid-list-bitmap
PolSci	000000111100
CS	111111000000
Math	000000000011



---

# Why Bitmap Index

---

- ❖ WHERE semid = 8 AND  
dept = 'CS'
- ❖ AND the two bitmaps to get  
000111000000
- ❖ Then get the rows
- ❖ Quick filtering on multiple not-  
so-unique columns
- ❖ Can use B+Tree, Hash Index

semid		rid-list-bitmap
-----	+	-----
10		111000111000
8		000111000100
1		000000000010
2		000000000001

dept		rid-list-bitmap
-----	+	-----
PolSci		000000111100
CS		111111000000
Math		000000000011



---

# Other Indexing Techniques

---

- ❖ So far: Indexing = ToC or “Back of Book” kind
- ❖ Indexing can be generalized to mean any performance enhancing storage technique
  - ❖ implies automatic maintenance in face of updates
- ❖ E.g.
  - ❖ Teradata® Join Index - an index to speed up joins
  - ❖ Postgres Partial Index - on a constrained table



---

# Relational Concepts and Indexing

---

- ❖ Candidate Key = unique identifier of row
- ❖ Primary Key = chosen candidate key
- ❖ uniqueness constraint = way to identify other candidate keys



---

# Primary Key vs Index

---

- ❖ Primary Index is the way the table records are stored
  - ❖ need not be unique, but needs enough distinct values
- ❖ Primary Key makes for good PI for most entity tables
  - ❖ If a column different from PK is being used often, it could make for a good PI candidate
- ❖ For table representing “relations” i.e. with multi-column primary keys, it’s usually best to have a non-unique primary index on the most common joining FK



---

# Secondary Index

---

- ❖ create [unique] index on <table>(<columns>)
- ❖ Unique constraints = unique secondary index
- ❖ Non-unique secondary index: for fast equality constraints on columns
  - ❖ Can be used for some kinds of aggregation operations
  - ❖ Can be ANDed with other SI's if there are equality constraints on different columns (not uncommon)
    - ❖ bitmap representation (bitmap index)