

# DB Algorithms: Accessing Data



# Basic Architecture

- ❖ ~~Application Layer - what most users see, talks SQL~~
- ❖ Parsing / Planning Layers - the intelligence
- ❖ Runtime or execution Layer - the brawn - NEXT
- ❖ ~~Storage Layer - where data resides, may include simple access layer~~

Applications

Parsing

Planning

Processing

Data Access

Data in SSD / HDD



---

# What do we have?

---

- ❖ Tables stored mostly as B+ Tree with a primary index
  - ❖ May be stored as Heap File with no primary index
  - ❖ May be stored using an Extendible Hashing structure
  - ❖ May be stored in columnar or hybrid format
  - ❖ May have multiple copies stored in different storage methods
- ❖ Secondary Indexes (Unique and Non-unique) - look like tables internally
  - ❖ This is an additional “sub-table” or structure
  - ❖ B+ Tree or Hash structure
- ❖ (Aggregate) Join Indexes, Materialized Views - may look like tables internally



# Simplest Query

```
db1=# select * from students;
```

sid	last_name	first_name	status
2016001	Shatdal	Ambuj	5
2015001	Lincoln	Abraham	1
2014101	Obama	Barack	2
2012144	Bush	George	4
2012101	Washington	George	3

(5 rows)

- ❖ Simply scan the table students to get all rows all columns

```
db1=# explain select * from students;  
          QUERY PLAN
```

```
-----  
Seq Scan on students (cost=0.00..1.05 rows=5 width=122)  
(1 row)
```



---

# WHERE clause

---

```
db1=# select * from students where last_name = 'Obama';
   sid   | last_name | first_name | status
-----+-----+-----+-----
 2014101 | Obama    | Barack    |      2
(1 row)
```

- ❖ Scan the table students to get rows
- ❖ Apply condition and keep the rows for which condition is true

```
db1=# explain select * from students where last_name = 'Obama';
               QUERY PLAN
-----
Seq Scan on students  (cost=0.00..1.06 rows=1 width=122)
  Filter: ((last_name)::text = 'Obama'::text)
(2 rows)
```



---

# WHERE clause

---

- ❖ No matter how complex the condition (excluding subqueries) on a single table, it can always be evaluated by looking at all rows
- ❖ This is the default method
- ❖ Subqueries are a different beast altogether



---

# Equality Conditions

---

- ❖ WHERE c1 = <value1> [... and c2 = <value 2> ...]
- ❖ 4 possibilities for the column(s)
  - ❖ primary index column(s) are a subset of the column(s)
  - ❖ There is a unique SI on a subset of the column(s)
  - ❖ There is a non unique SI on a subset of the column(s)
  - ❖ None of the above



---

# Equality Condition - PI Access Path

---

- ❖ If there is a PI available, one could go through the primary index to find the rows
  - ❖ Look up the value in the B+ Tree index
  - ❖ Find the leaf node
  - ❖ Find the row(s) containing the value
  - ❖ apply any “residual conditions”
  - ❖ Return the qualifying row(s)



---

# Residual Conditions

---

- ❖ Conditions other than the ones that apply to the index being used



---

# PI Access Path Example

---

```
db1=# select * from courses where dept = 'CS' and cid = 564;
```

dept	cid	name
CS	564	DBMS

(1 row)

- ❖ If there WHERE has equality on the primary index (primary key) then we could also use the index to locate the row(s)

```
db1=# explain select * from courses where dept='CS' and cid = 564;
```

QUERY PLAN

---

```
Index Scan using courses_pkey on courses (cost=0.15..8.17 rows=1 width=98)
  Index Cond: (((dept)::text = 'CS'::text) AND (cid = 564))
(2 rows)
```



---

# Equality Condition - USI Access Path

---

- ❖ If there is a USI available, one could go through the secondary index to find the rows
  - ❖ Look up the value in the index
  - ❖ Find the leaf / data node
  - ❖ Find the *rid* of the row containing the value
  - ❖ Look up the row using the *rid*
  - ❖ Apply any residual conditions
  - ❖ Return the qualifying row, if any



---

# Equality Condition - NUSI Access Path

---

- ❖ If there is a NUSI available, one could go through the secondary index to find the rows
  - ❖ Look up the value in the index
  - ❖ Find the leaf/ data node
  - ❖ Find the *rids* of the rows containing the value
    - ❖ this could be a bitmap, for example
  - ❖ Sort the *rids*
  - ❖ Look up the rows using the *rids*
  - ❖ Apply any residual conditions
  - ❖ Return the qualifying row, if any



# NUSI Access Path Example

```
hw2=# select * from sales where dept = 1;
 store | dept | weekdate | weeklysales
-----+-----+-----+-----
      1 |    1 | 2010-02-05 |      24924.5
      1 |    1 | 2010-02-12 |      46039.5
etc...
```

- ❖ If there WHERE has equality on the primary index (primary key) then we could also use the index to locate the row(s)

```
hw2=# explain select * from sales where dept = 1;
               QUERY PLAN
```

```
-----
Bitmap Heap Scan on sales  (cost=129.24..2493.43 rows=6815 width=16)
```

```
  Recheck Cond: (dept = 1)
```

```
    -> Bitmap Index Scan on salesdept  (cost=0.00..127.53 rows=6815 width=0)
```

```
          Index Cond: (dept = 1)
```

```
(4 rows)
```



# No Index - Scan and Evaluate

```
db1=# select * from students where last_name = 'Obama';
   sid   | last_name | first_name | status
-----+-----+-----+-----
 2014101 | Obama    | Barack    |      2
(1 row)
```

- ❖ Scan the table students to get rows
- ❖ Apply condition and keep the rows for which condition is true

```
db1=# explain select * from students where last_name = 'Obama';
               QUERY PLAN
-----
Seq Scan on students  (cost=0.00..1.06 rows=1 width=122)
  Filter: ((last_name)::text = 'Obama'::text)
(2 rows)
```



---

# Multiple NUSI's

---

- ❖ WHERE  $C1 = 'A'$  and  $C2 = 10$
- ❖ Both are non-unique and both have a NUSI
- ❖ Lookup both the indexes
- ❖ AND the *rid*-lists (bitmap or otherwise)
- ❖ Sort the *rids*
- ❖ Look up the rows and apply residual conditions



---

# SQL Operations

---

- ❖ SELECT query blocks
- ❖ Optional UNION, INTERSECT, EXCEPT operations on them



---

# SQL Operations - SELECT block

---

- ❖ Select rows from one or more tables
- ❖ If more than 1 table, must join (or cross join) them
  - ❖ Can have fairly involved WHERE conditions
- ❖ Optionally aggregate the rows - GROUP BY
  - ❖ Evaluate Having - can get complex
- ❖ Optionally evaluate ordered-analytic (window) functions
- ❖ Optionally do a DISTINCT