

Query Planning: Part 1

Basic Architecture

- ❖ ~~Application Layer - what most users see, talks SQL~~
- ❖ Parsing / Planning Layers - the intelligence - NEXT
- ❖ ~~Runtime or execution Layer - the brawn~~
- ❖ ~~Storage Layer - where data resides, may include simple access layer~~

Applications

Parsing

Planning

Processing

Data Access

Data in SSD / HDD


Simple Example

❖ `SELECT last_name, first_name
FROM presidents
WHERE num = 42;`

Step 1: Parsing

```
SELECT last_name, first_name  
FROM presidents  
WHERE num = 42;
```

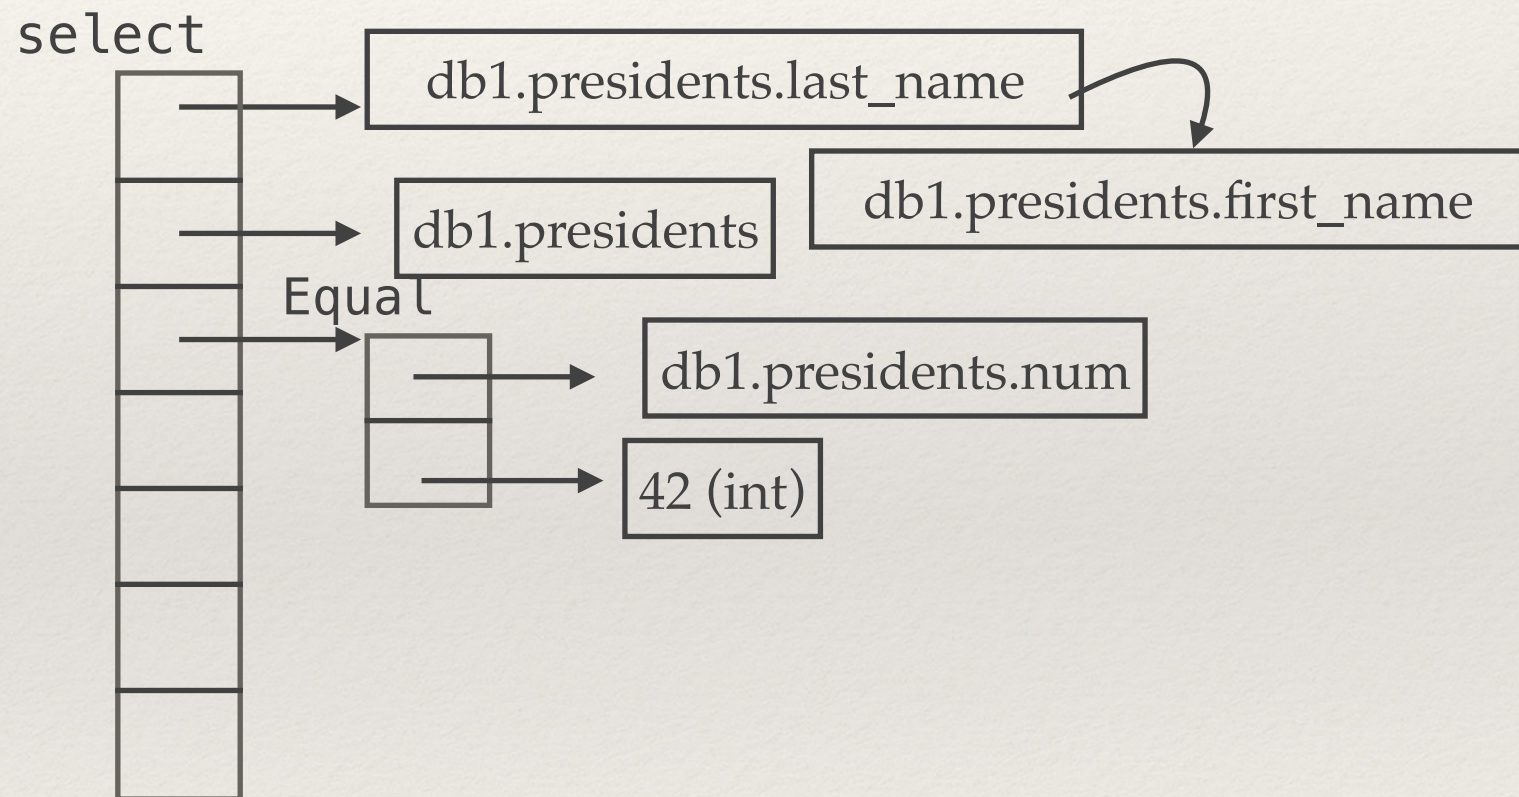
- ❖ The SQL is given to a parser that makes a parse tree
- ❖ simple syntax errors



Step 2: Resolving the references

- ❖ The references are looked up in the catalog a.k.a. (data dictionary)
- ❖ errors if things are not found
- ❖ E.g.
 - ❖ presidents table/view must exist
 - ❖ last_name, first_name, num have to be columns from presidents table/view

```
SELECT last_name, first_name  
FROM presidents  
WHERE num = 42;
```



System Catalog

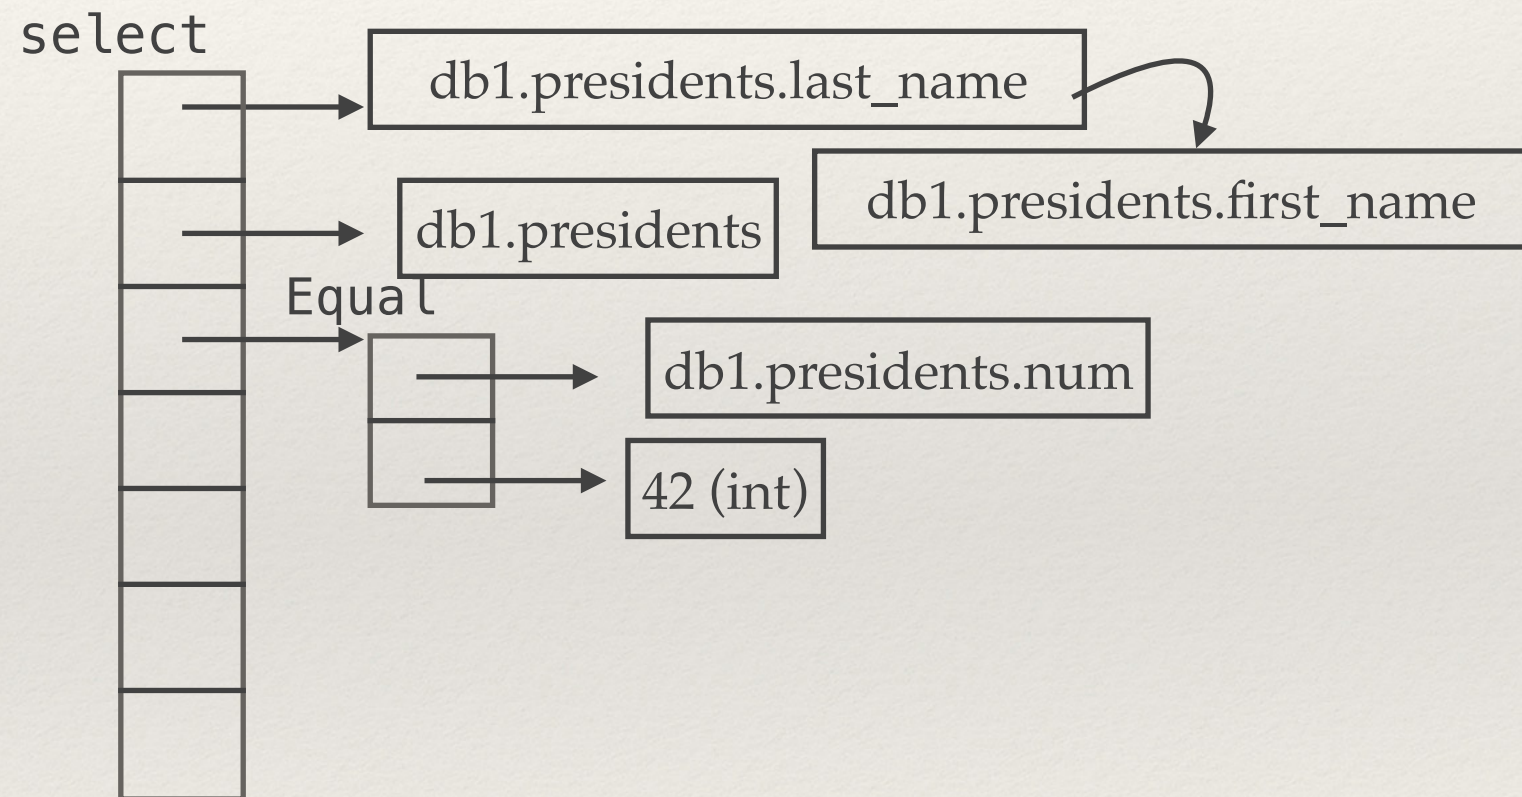

- ❖ Set of tables that maintain metadata
 - ❖ Users / Databases / Schemas
 - ❖ Tables, Views (names, # columns, view_def / pointer to file)
 - ❖ Indexes (name, table(s), columns, kind)
 - ❖ Attributes (name, type, nullability)
 - ❖ Access Rights
 - ❖ Statistics
 - ❖ User defined functions
 - ❖

Resolution Process

- ❖ The Planner has to issue (highly optimized) queries against the catalog to get all the needed information
- ❖ Multiple iterations (e.g. once we get a view, we need to get all the view references resolved)
- ❖ At the end: no views etc.
 - ❖ all references are to base tables and their columns

SQL-ParseTree-RA Operator Tree

```
SELECT last_name, first_name  
FROM presidents  
WHERE num = 42;
```



The Optimization Process

- ❖ Goal: To generate a query execution plan that is as optimal as practicable
- ❖ Both rules-based and cost-based
- ❖ Rules-based - where the more optimal option is clear
- ❖ Cost-based - where the options have to be compared
- ❖ Think of it as transformation of “original RA” tree to an augmented (and hopefully optimal)RA tree

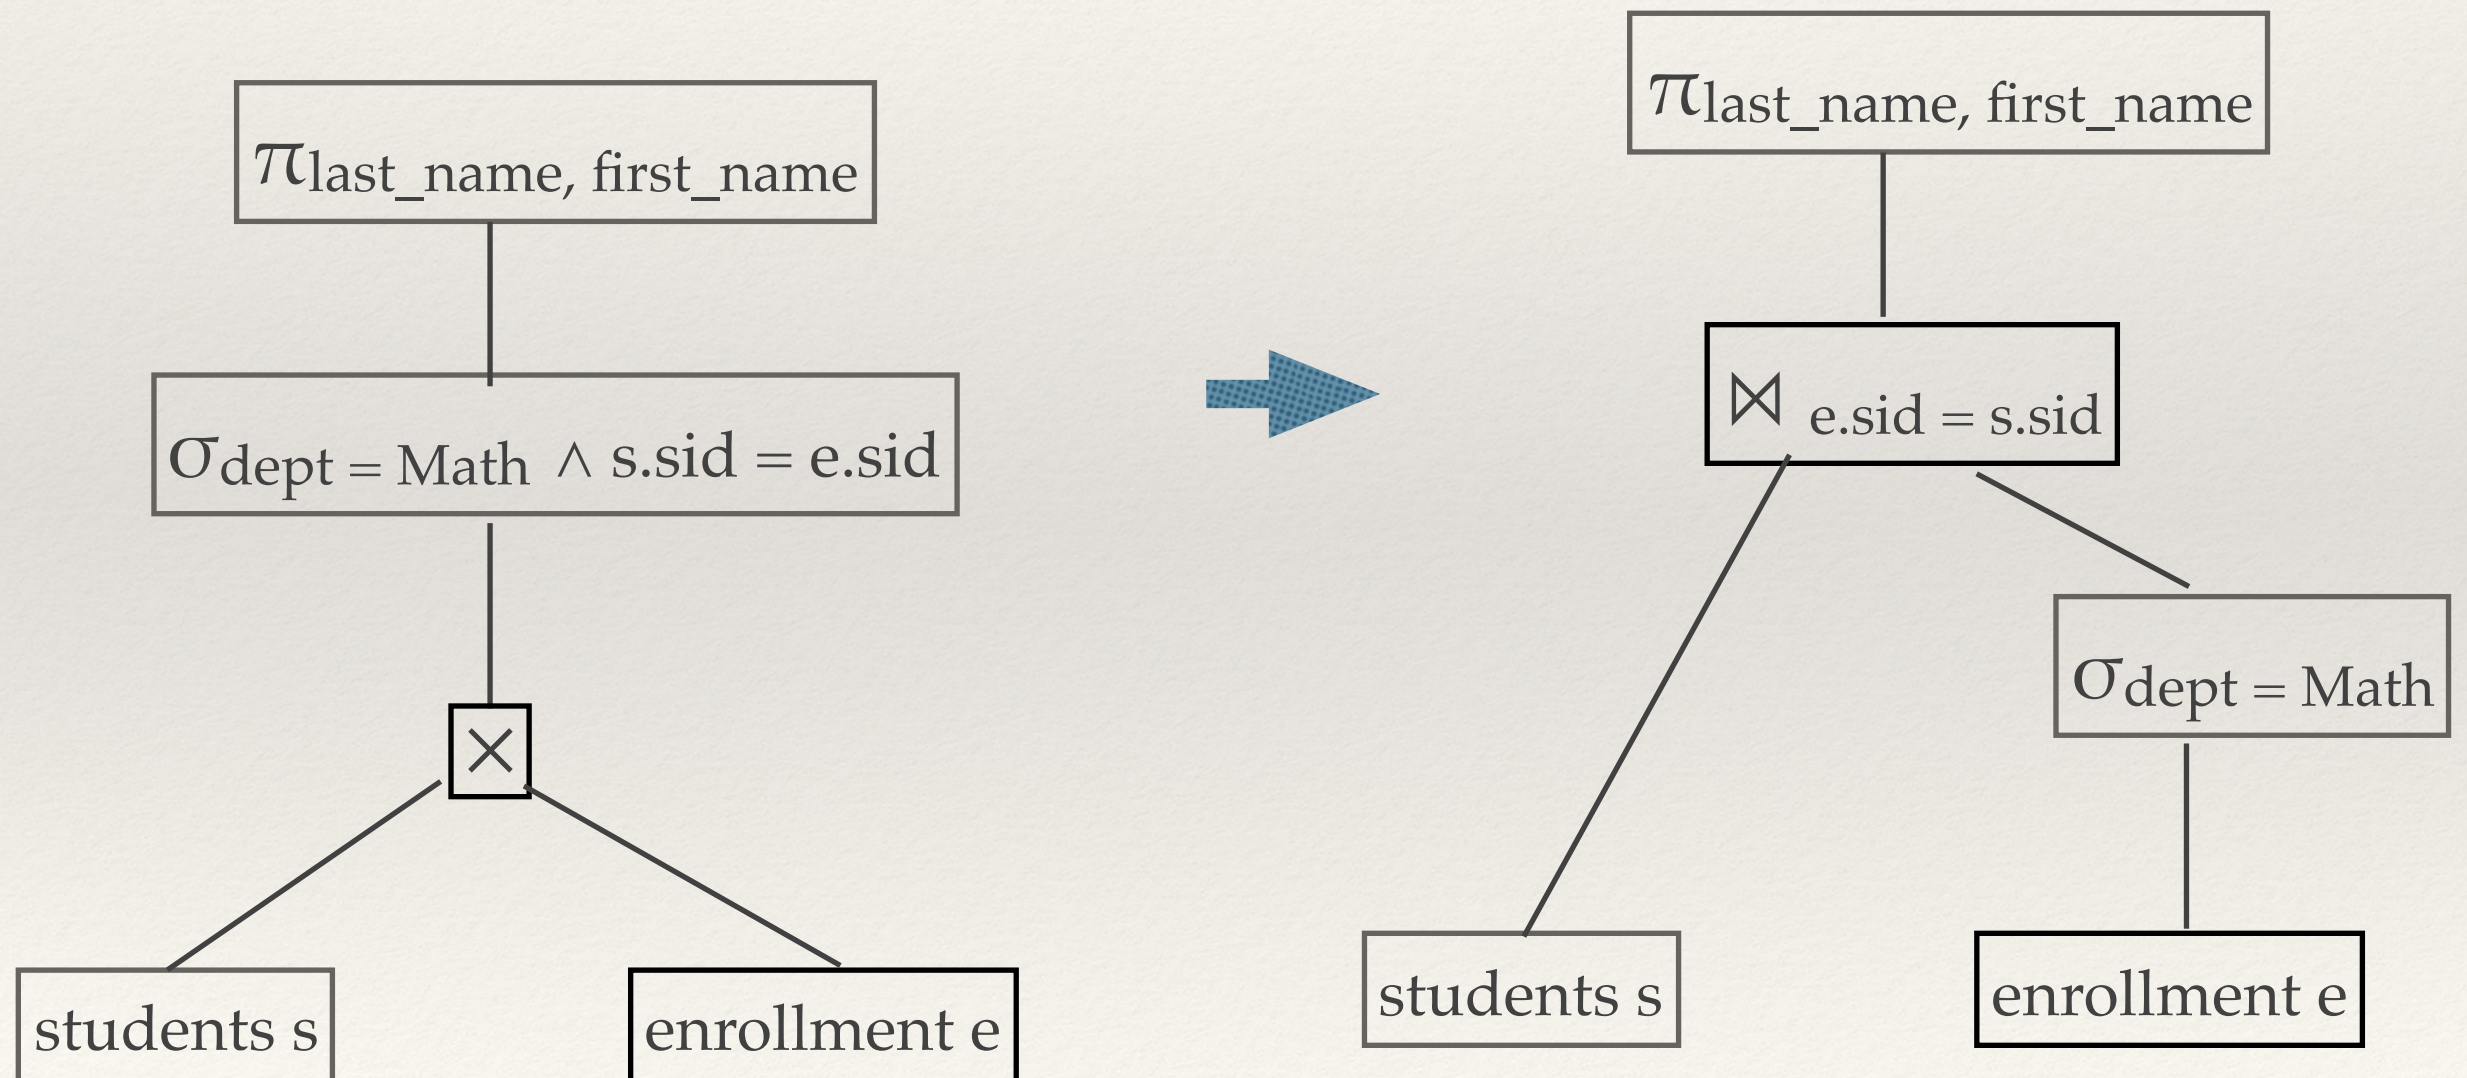
Basic SPJ Query

- ❖ This query has the three basic relational algebra operations
- ❖ select i.e. condition
- ❖ project i.e. choosing a subset of columns and keeping distinct values
- ❖ join (equi-join) on FK-PK relationship

```
SELECT DISTINCT last_name, first_name
FROM students s, enrollment e
WHERE s.sid = e.sid
AND e.dept = 'Math';
```

last_name	first_name
Bush	George

An SPJ Query



Query Optimization

- ❖ Goal: to find equivalent relational operator tree that would produce the desired result using least resources and / or most quickly
- ❖ If the equivalent transformation is “obvious” then we can create a rule, e.g. “push selections to the table”
- ❖ If the equivalent transformations are not obvious we need to compare costs of different ways of doing things

Cost Estimation

- ❖ Cost: number of I/O's + CPU + messaging
 - ❖ I/O costs dominate
 - ❖ All proportional to size of input being processed
 - ❖ Row size (easy!) x number of rows (hard)
- ❖ Estimate I/O cost for performing operation given the input estimates, but CPU cost are usually estimated too
- ❖ Then produce the output size estimate (cardinality + row size) and whether it's sorted (and the sort key)
- ❖ Pipelining can affect the I/O cost

Result Size Estimation - Single Table

- ❖ Est. Result Size = # Rows * Estimated Reduction Factor
- ❖ WHERE column = value
 - ❖ Estimate: #Rows * (1 / #unique values)
 - ❖ #unique values estimated from statistics
 - ❖ statistics on the column would normally keep number of unique values, min, max and possibly a histogram
 - ❖ could sample a page
 - ❖ else default to something, say #Rows * (1 / 10)

Result Size Estimation – contd.

- ❖ WHERE column > value
 - ❖ Estimate: $\# \text{Rows} * [(\text{max} - \text{value}) / (\text{max} - \text{min})]$
 - ❖ needs with basic statistics or histogram
- ❖ WHERE column IN (v1, v2)
 - ❖ same as column = v1 OR column = v2
 - ❖ Add up the equality estimates
 - ❖ heuristic: limit to $\# \text{Rows} / 2$

Data Distribution


- ❖ Basic Statistics

- ❖ $\text{min} = 0$

- ❖ $\text{max} = 14$

- ❖ $\#\text{unique} = 14$

- ❖ $\#\text{rows} = 45$



Using Statistics

- ❖ Basic Statistics

- ❖ `min = 0`
- ❖ `max = 14`
- ❖ `#unique = 14`
- ❖ `#rows = 45`

- ❖ Estimate # rows

- ❖ `value = 7`
- ❖ `value > 12`
- ❖ `value in (5,11)`

(Equi-width) Histogram

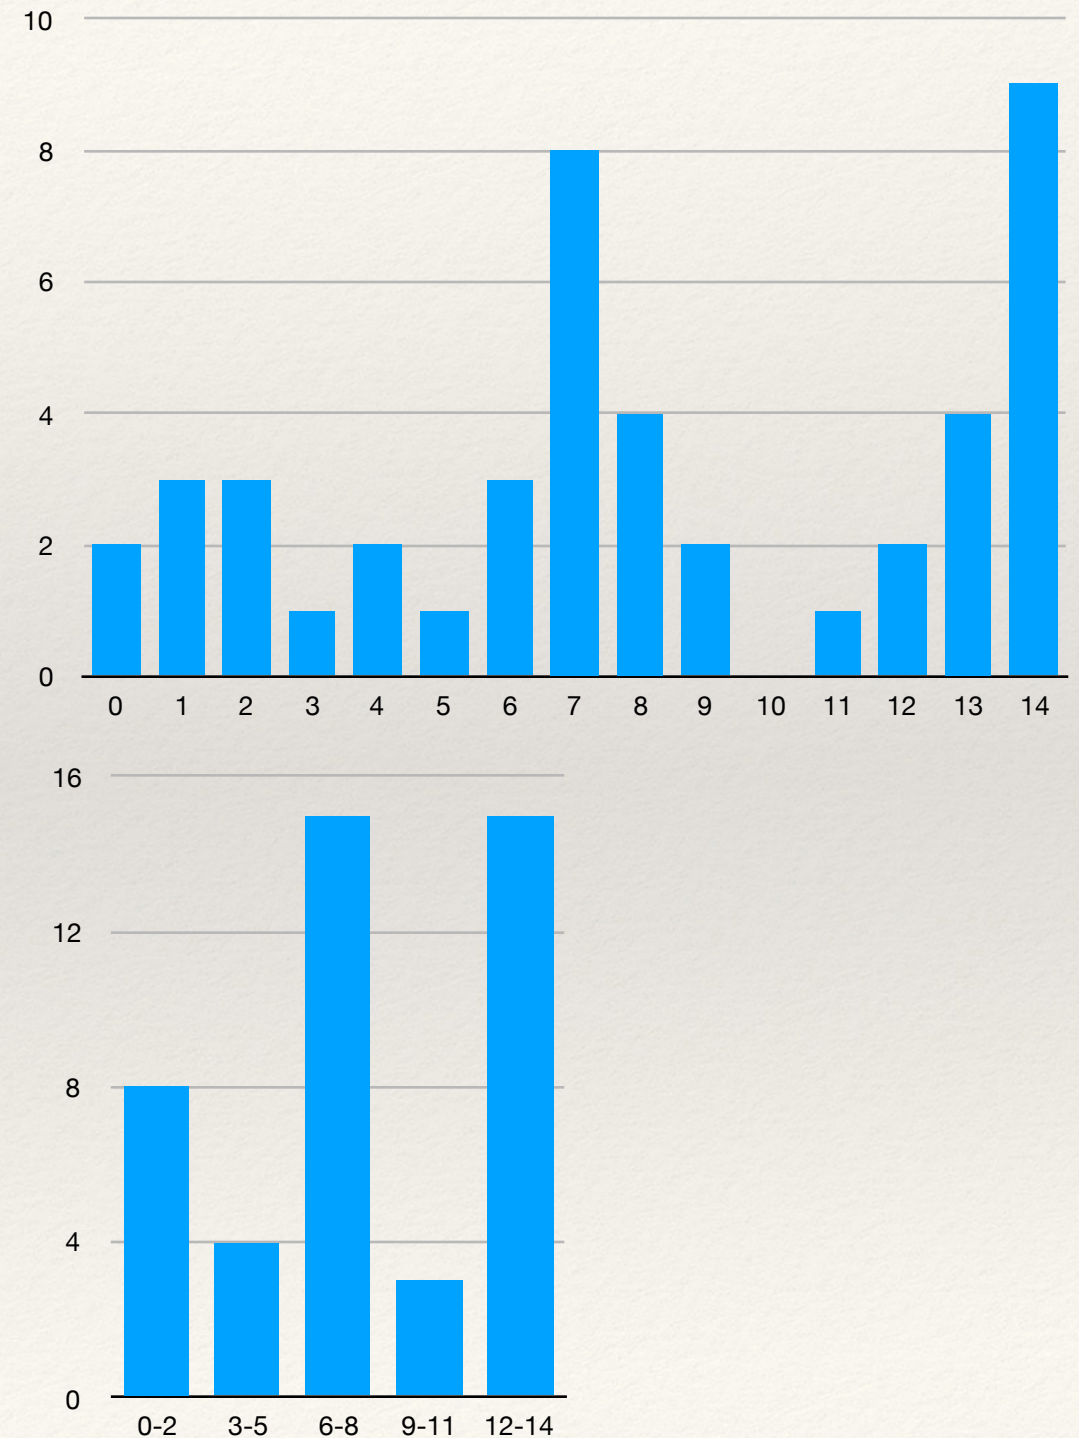
❖ Basic Statistics

❖ $\text{min} = 0$

❖ $\text{max} = 14$

❖ $\#\text{unique} = 14$

❖ $\#\text{rows} = 45$



(Equi-height) Histogram



❖ Basic Statistics

❖ $\text{min} = 0$

❖ $\text{max} = 14$

❖ $\#\text{unique} = 14$

❖ $\#\text{rows} = 45$



Join Size Estimation

- ❖ Cross Join of R, S, $\#rows = \#rows(R) * \#rows(S)$
- ❖ WHERE R.column1 = S.column2
 - ❖ Assuming PK-FK join
 - ❖ i.e. joining columns are a key to one of the tables
 - ❖ $\#rows \text{ in result} = \min(\#rows(R), \#rows(S))$

Join Size Estimation - General

- ❖ Assume A is the joining attribute
 - ❖ $n_A(R)$ and $n_A(S)$ are number of unique values of A in R and S respectively
- ❖ For a row r in R , it would join with $\#rows(S) / n_A(S)$
 - ❖ \Rightarrow total row estimate is $\#rows(R) * \#rows(S) / n_A(S)$
 - ❖ By symmetry it's also $\#rows(R) * \#rows(S) / n_A(R)$
- ❖ Normally we expect $n_A(S) = n_A(R)$, but if there are non-joining values, one would use the smaller of the two estimates
- ❖ Note that this gives the same estimate for FK-PK case
- ❖ Non-equijoins are estimated cross joins followed by estimate for the individual conditions

Join Size Estimation and Histograms

- ❖ Histograms make the estimates more accurate

$$Estimate = \sum_{bucket=1}^n n(R) * n(S) / \max(n_A(R), n_A(S))$$

- ❖ No Histogram Estimate = 144, Actual Value = 223

Equi-width Histogram Estimate

Bucket	Freq	Estimate
0-2	8	21.3
3-5	4	5.3
6-8	15	75.0
9-11	3	3.0
12-14	15	75.0
	45	179.7

Equi-depth Histogram Estimate

Bucket	Freq	Estimate
0-3	9	20.3
4-6	6	12.0
7-7	8	64.0
8-13	13	28.2
14-14	9	81.0
	45	205.4

Other Estimates

- ❖ Projection / Group By estimate
 - ❖ Very hard unless there is direct statistics available for the group by / projected columns, sampling is of limited help
- ❖ Set operations - fairly approximate
 - ❖ UNION: $\#rows(R) + \#rows(S)$
 - ❖ INTERSECT: $\min(\#rows(R), \#rows(S))$
 - ❖ EXCEPT: $\#rows(R)$
- ❖ Outer joins - upper bounds
 - ❖ R left outer join S: $\text{estimate}(R \text{ join } S) + \#rows(R)$, ROJ is symmetric
 - ❖ FULL OJ: $\text{estimate}(R \text{ join } S) + \#rows(R) + \#rows(S)$

Histograms

- ❖ Besides nRows, min, max, unique count, capture the distribution of data
- ❖ For each bucket
 - ❖ frequency for range of values
 - ❖ number of unique values in bucket
- ❖ equi-depth vs. equi-width
- ❖ equi-depth usually captures more information by keeping more details about frequent values
- ❖ Some databases may even keep separate info for modal value

Transformations

- ❖ Two expressions are equivalent if and only if they produce the same result on every legal instance
- ❖ The goal is to find the least costly but equivalent expression for evaluating the query
- ❖ Most are intuitive

Selections & Projections

- ❖ $\sigma_{c1 \wedge c2 \wedge c3}(R) = \sigma_{c1}(\sigma_{c2}(\sigma_{c3}(R)))$
- ❖ since \wedge is commutative $\sigma_{c1 \wedge c2}(R) = \sigma_{c2 \wedge c1}(R)$
 - ❖ $\Rightarrow \sigma_{c1}(\sigma_{c2}(R)) = \sigma_{c2}(\sigma_{c1}(R))$
- ❖ reason why we convert conditions to CNF
- ❖ A sequence of projections can be reduced to final one
 - ❖ $\pi_{L1}(\pi_{L2}(\pi_{L3}(R))) = \pi_{L1}(R)$

Joins and Cross Products

- ❖ Cross product and natural (equi-) join are commutative
 $R \times S = S \times R$ and $R \bowtie S = S \bowtie R$
- ❖ They are also associative
 $R \times (S \times T) = (R \times S) \times T$ and $R \bowtie (S \bowtie T) = (R \bowtie S) \bowtie T$
- ❖ Together, we can join them in any order we choose

Other SPJ Equivalences

- ❖ $\sigma_{c1}(\pi_{L1}(R)) = \pi_{L1}(\sigma_{c1}(R))$ if selection only needs attributes retained by projection
- ❖ $\sigma_c(R \bowtie S) = \sigma_c(R) \bowtie S$ if c contains attributes from R only
- ❖ Generalizing the above - pushing conditions through joins
 - ❖ $\sigma_{c1 \wedge c2 \wedge c3}(R \bowtie S) = \sigma_{c1}(\sigma_{c2}(R) \bowtie \sigma_{\overline{c2}}(S))$
c3
 - ❖ where $c2$ contains attributes only from R
 - ❖ $c3$ contains attributes only from S
 - ❖ $c1$ contains attributes from both

Other SPJ Equivalences

- ❖ Pushing Projections

- ❖ $\pi_L(R \bowtie_C S) = \pi_{L1}(R) \bowtie_C \pi_{L2}(S)$ where

- ❖ L1 is subset of attributes of R that are in L

- ❖ L2 is subset of attributes of S that are in L

- ❖ All attributes in C must be there in L

- ❖ Generalizes to

- ❖ $\pi_L(R \bowtie_C S) = \pi_L(\pi_{L1}(R) \bowtie_C \pi_{L2}(S))$ where

- ❖ L1 is subset of attributes of R that are in L or C

- ❖ L2 is subset of attributes in S that are in L or C

Basic Optimization Strategy

- ❖ Enumerate alternative plans
- ❖ Pick the best (least estimated cost)
- ❖ Equivalences are the basis for alternate plans. This is combined with alternative operator algorithms and alternate ways to access a single table

Single Table Access

- ❖ Enumerate all ways to evaluate a single table access
 - ❖ Scan and evaluate condition
 - ❖ Use the primary (clustered) index if applicable
 - ❖ Use one secondary index if applicable
 - ❖ Use multiple secondary indices if applicable
 - ❖ Use other structures (e.g. materialized views) if applicable

Joins

- ❖ left - right asymmetry in hash and nested loop
- ❖ Assume hash and sort-merge
- ❖ 2 table joins $\Rightarrow R \bowtie S$ (hash), $S \bowtie R$ (hash), $R \bowtie S$ (SM)
- ❖ 3 tables - just one algorithm, ignore asymmetry
 - ❖ $R \bowtie (S \bowtie T)$, $(R \bowtie S) \bowtie T$, $(R \times T) \bowtie S$