

Transaction Management Overview

Chapter 16

Transactions

- ❖ Concurrent execution of user programs is essential for good DBMS performance.
 - Hides waits for disk I/O by keeping several transaction going at once
 - Exploit multi-core architecture
- ❖ DBMS is only concerned about what data is read/written from/to the database during the execution of a user program
- ❖ A transaction is the DBMS's abstract view of a user program: a sequence of reads and writes.

Concurrency in a DBMS

- ❖ Users submit transactions, and can think of each transaction as executing by itself.
 - Concurrency is achieved by the DBMS, which interleaves actions (reads/writes of DB objects) of various transactions.
 - Transactions are required to leave DBMS in a consistent state, i.e. each transaction must leave the database in a consistent state if the DB is consistent when the transaction begins.
- ❖ Issues: Effect of *interleaving* transactions, and *crashes*.

Atomicity of Transactions

- ❖ transaction might *commit* after completing all its actions, or it could *abort* (or be aborted by the DBMS) after executing some actions.
- ❖ A very important property guaranteed by the DBMS for all transactions is that they are *atomic*. That is, a user can think of a Xact as always executing all its actions in one step, or not executing any actions at all.

Example

- ❖ Consider two transactions (*Xacts*):

| | | | | |
|-----|-------|-----------|----------|-----|
| T1: | BEGIN | A=A+100, | B=B-100 | END |
| T2: | BEGIN | A=1.06*A, | B=1.06*B | END |

- ❖ Intuitively, the first transaction is transferring \$100 from B's account to A's account. The second is crediting both accounts with a 6% interest payment.
- ❖ There is no guarantee that T1 will execute before T2 or vice-versa, if both are submitted together. However, the net effect *must* be equivalent to these two transactions running serially in some order.

Example (Contd.)

- ❖ Consider a possible interleaving (schedule):

| | | |
|-----|-----------|----------|
| T1: | A=A+100, | B=B-100 |
| T2: | A=1.06*A, | B=1.06*B |

- ❖ This is OK. But what about:

| | | |
|-----|--------------------|---------|
| T1: | A=A+100, | B=B-100 |
| T2: | A=1.06*A, B=1.06*B | |

- ❖ The DBMS's view of the second schedule:

| | | |
|-----|------------------------|------------|
| T1: | R(A), W(A), | R(B), W(B) |
| T2: | R(A), W(A), R(B), W(B) | |

Scheduling Transactions

- ❖ Serial schedule: Schedule that does not interleave the actions of different transactions.
- ❖ Equivalent schedules: For any database state, the effect (on the set of objects in the database) of executing the first schedule is identical to the effect of executing the second schedule.
- ❖ Serializable schedule: A schedule that is equivalent to some serial execution of the transactions.

(Note: If each transaction preserves consistency, every serializable schedule preserves consistency.)

Anomalies with Interleaved Execution

- ❖ Reading Uncommitted Data (WR Conflicts, “dirty reads”):

| | | |
|-----|---------------|-------------------|
| T1: | R(A), W(A), | R(B), W(B), Abort |
| T2: | R(A), W(A), C | |

- ❖ Unrepeatable Reads (RW Conflicts):

| | | |
|-----|---------------|---------------|
| T1: | R(A), | R(A), W(A), C |
| T2: | R(A), W(A), C | |

Anomalies (Continued)

❖ Overwriting Uncommitted Data (WW Conflicts):

| | | |
|-----|---------------|---------|
| T1: | W(A), | W(B), C |
| T2: | W(A), W(B), C | |

Lock-Based Concurrency Control

- ❖ Strict Two-phase Locking (Strict 2PL) Protocol:
 - Each Xact must obtain a **S (shared) lock** on object before reading, and an **X (exclusive) lock** on object before writing.
 - All locks held by a transaction are released when the transaction completes (i.e. at commit/abort).
 - **(Non-strict) 2PL Variant:** Release locks anytime, but cannot acquire locks after releasing any lock.
 - If an Xact holds an X lock on an object, no other Xact can get a lock (S or X) on that object.
- ❖ Strict 2PL allows only serializable schedules.
 - Additionally, it simplifies transaction aborts
 - **(Non-strict) 2PL** also allows only serializable schedules, but involves aborts may have to be cascaded to transactions that read the data modified by this now-aborted Xact

Aborting a Transaction

- ❖ If a transaction T_i is aborted, all its actions have to be undone. Not only that, if T_j reads an object last written by T_i , T_j must be aborted as well (possible with non-strict 2PL)!
- ❖ Most systems avoid such *cascading aborts* by releasing a transaction's locks only at commit time (strict 2PL).
 - If T_i writes an object, T_j can read this only after T_i commits.
- ❖ In order to *undo* the actions of an aborted transaction, the DBMS maintains a *log* in which every write is recorded. This mechanism is also used to recover from system crashes: all active Xacts at the time of the crash are aborted when the system comes back up.

The Log

- ❖ The following actions are recorded in the log:
 - *Ti writes an object:* the old value and the new value.
 - Log record must go to disk before the changed page!
 - *Ti commits/aborts:* a log record indicating this action.
- ❖ Log records are chained together by Xact id, so it's easy to undo a specific Xact.
- ❖ Log is often *duplexed* and *archived* on stable storage.
- ❖ All log related activities (and in fact, all CC related activities such as lock/unlock, dealing with deadlocks etc.) are handled transparently by the DBMS.

Summary

- ❖ Concurrency control and recovery are among the most important functions provided by a DBMS.
- ❖ Users need not worry about concurrency.
 - System automatically inserts lock/unlock requests and schedules actions of different Xacts in such a way as to ensure that the resulting execution is equivalent to executing the Xacts one after the other in some order.
- ❖ Write-ahead logging (WAL) is used to undo the actions of aborted transactions and to restore the system to a consistent state after a crash.
 - *Consistent state*: Only the effects of committed Xacts seen.