

Transactions in SQL

Transactions

- ❖ If nothing is specified, automatically for each individual statement
- ❖ Begin Transaction (or Start Transaction)
- ❖ followed by the statements that are part of that transaction
- ❖ Commit or Rollback

Example

```
db1=# begin transaction;
```

```
BEGIN
```

```
db1=# select * from sp;
```

name	item	price
S2	P3	100
S1	P2	20
S1	P1	10
S3	P4	1000
S2	P1	11
S4	P1	9
	P1	
S5		
S4	P3	

(9 rows)

Update Table

```
db1=# update sp set price = price * 2;
```

```
UPDATE 9
```

```
db1=# select * from sp;
```

name	item	price
S2	P3	200
S1	P2	40
S1	P1	20
S3	P4	2000
S2	P1	22
S4	P1	18
	P1	
S5		
S4	P3	

(9 rows)

Change your mind - rollback

```
db1=# rollback;
```

```
ROLLBACK
```

```
db1=# select * from sp;
```

name	item	price
S2	P3	100
S1	P2	20
S1	P1	10
S3	P4	1000
S2	P1	11
S4	P1	9
	P1	
S5		
S4	P3	

(9 rows)

Savepoint

- ❖ In longer transactions, mark places to rollback to if necessary

```
db1=# begin transaction;  
BEGIN  
db1=# update sp set price = price * 2;  
UPDATE 9
```

```
db1=# select * from sp;
```

name	item	price
S2	P3	200
S1	P2	40
S1	P1	20
S3	P4	2000
S2	P1	22
S4	P1	18
	P1	
S5		
S4	P3	

(9 rows)

```
db1=# savepoint save1;  
SAVEPOINT
```


Savepoint

- ❖ rollback to a previously saved (named) save point

```
db1=# update sp set price = price * 3;  
UPDATE 9
```

```
db1=# select * from sp;
```

name	item	price
S2	P3	600
S1	P2	120
S1	P1	60
S3	P4	6000
S2	P1	66
S4	P1	54
	P1	
S5		
S4	P3	

(9 rows)

```
db1=# rollback to savepoint save1;  
ROLLBACK
```


Commit

- ❖ Commit will end a transaction keeping the final state

```
db1=# select * from sp;  
name | item | price
```

-----+	-----+	-----
S2	P3	200
S1	P2	40
S1	P1	20
S3	P4	2000
S2	P1	22
S4	P1	18
	P1	
S5		
S4	P3	

```
(9 rows)
```

```
db1=# commit;  
COMMIT
```

What to Lock?

- ❖ DB entities are nested in hierarchies, e.g.
 - ❖ Table
 - ❖ Page
 - ❖ Row
- ❖ Or a B+ Tree hierarchy
 - ❖ Root
 - ❖ Subtree at any point
 - ❖ Leaf node
 - ❖ Row / Record

General Rules

- ❖ Lock what will be accessed / modified
- ❖ Lower granularity allows higher concurrency
- ❖ Lower granularity locks may allow “phantom” problem.
 - ❖ A concurrent transaction may insert new rows and commit.
 - ❖ Next time the same data is read, it would see the newly inserted row

Isolation Levels

- ❖ DBMS and SQL in practice offer varying levels of consistency
 - ❖ Serializable
 - ❖ Repeatable Read
 - ❖ Read Committed
 - ❖ Read Uncommitted

Serializable

- ❖ The highest level of consistency
- ❖ 2 Phase locking, ensure no phantoms by either table level or B+tree based locking (to prevent addition to locked pages)

Repeatable Read

- ❖ In practice, could be same as serializable with 2 Phase locking, but don't guarantee the phantom problem
- ❖ Essentially lock every read item but no need to lock “higher” level objects to prevent phantoms

Read Committed

- ❖ Allows others transactions committed data to be visible while the transaction is in flight
 - ❖ Its own changes are not visible to others
- ❖ Short read (shared) locks that are released immediately.
- ❖ Exclusive locks are released at commit / rollback(abort).
- ❖ Good practical level for long running read-mostly transactions - e.g. data analysis where you may even want to see data that's recently added

Read Uncommitted

- ❖ Essentially equivalent to no read / shared locks
- ❖ Can read dirty (uncommitted) data from other transactions
- ❖ Exclusive locks are acquired and kept till commit
- ❖ Not all systems support this (e.g. Postgres, Teradata).