

Relational Algebra: Part 1

Relational Algebra and Calculus

- ❖ Two mathematical Query Languages are the basis for SQL and its implementation
- ❖ Relational Algebra: More operational, useful for
 - ❖ representing execution plans, a foundation for SQL execution
- ❖ Relational Calculus: Lets users describe what they want, rather than how to compute it. (Non-operational, declarative.)
 - ❖ SQL is the usable result from Relational Calculus

Preliminaries

- ❖ A query is applied to relation instances (i.e. current “value” of the relation), and the result of a query is also a relation instance.
- ❖ Schemas of input relations for a query are fixed
- ❖ The schema for the result of a given query is also pre-determined by the query definition

Example Instances

R1

sid	bid	day
22	101	1996-10-10
58	103	1996-11-12

S1

sid	sname	rating	age
22	dustin	7	45
31	lubber	8	55.5
58	rusty	10	35

S2

sid	sname	rating	age
31	lubber	8	55.5
58	rusty	10	35
28	yuppy	9	35
44	guppy	5	35

Algebraic Operations

- ❖ Basic operations:
 - ❖ Selection (σ) Selects a subset of rows from relation.
 - ❖ Projection (π) Deletes unwanted columns from relation.
 - ❖ Cross-product (\times) Allows us to combine two relations.
 - ❖ Set-difference ($-$) Tuples in relation 1, but not in relation 2.
 - ❖ Union (\cup) Tuples in relation1 and in relation 2
- ❖ Additional operations:
 - ❖ Intersection(\cap), join(\Join), division($/$), renaming: not essential
- ❖ Since each operation returns a relation operations can be composed

Projection

- ❖ Deletes attributes that are not in projection list.
- ❖ Schema of result contains exactly the fields in the projection list, with the same names that they had in the input relation
- ❖ Projection operator has to eliminate duplicates.
 - ❖ In practice, systems don't do duplicate elimination unless asked

$\pi_{\text{sname}, \text{rating}}(\text{S2})$

sname	rating
lubber	8
rusty	10
yuppy	9
guppy	5

$\pi_{\text{age}}(\text{S2})$

age
55.5
35

sid	sname	rating	age
31	lubber	8	55.5
58	rusty	10	35
28	yuppy	9	35
44	guppy	5	35

Selection

- ❖ Selects rows that satisfy selection condition.
- ❖ *Schema* of result identical to schema of input relation.
- ❖ *Result* relation can be the *input* for another relational algebra operation

sid	sname	rating	age
31	lubber	8	55.5
58	rusty	10	35
28	yuppy	9	35
44	guppy	5	35

$\sigma_{\text{rating} > 8} (S2)$

sid	sname	rating	age
58	rusty	10	35
28	yuppy	9	35

$\pi_{\text{sname}, \text{rating}} (\sigma_{\text{rating} > 8} (S2))$

sname	rating
rusty	10
yuppy	9

Union/Intersection/Minus

- ❖ All of these operations take two input relations, which must be union-compatible:
 - ❖ Same number of fields.
 - ❖ Corresponding fields have the same type.

$S1 - S2$

sid	sname	rating	age
22	dustin	7	45

$S1 \cup S2$

sid	sname	rating	age
31	lubber	8	55.5
22	dustin	7	45
58	rusty	10	35
44	guppy	5	35
28	yuppy	9	35

$S1 \cap S2$

sid	sname	rating	age
31	lubber	8	55.5
58	rusty	10	35

Cross Product

- ❖ Each row of S1 is paired with each row of R1.
- ❖ Result schema has one field per field of S1 and R1 with field names *inherited* if possible.
 - ❖ Conflict: Both S1 and R1 have a field called sid.
 - ❖ Renaming operation helps here: $\rho(C(1 \rightarrow \text{sid1}, 5 \rightarrow \text{sid2}), S1 \times R1)$

sid	sname	rating	age	sid	bid	day
22	dustin	7	45	22	101	1996-10-10
31	lubber	8	55.5	22	101	1996-10-10
58	rusty	10	35	22	101	1996-10-10
22	dustin	7	45	58	103	1996-11-12
31	lubber	8	55.5	58	103	1996-11-12
58	rusty	10	35	58	103	1996-11-12

Joins

- ❖ Condition (θ) Join: $R \bowtie_{\theta} S = \sigma_{\theta}(R \times S)$
- ❖ Result schema is same as that of cross product

$S1 \bowtie_{(s1.sid < R1.sid)} R1$

sid	sname	rating	age	sid	bid	day
22	dustin	7	45	22	101	1996-10-10
31	lubber	8	55.5	22	101	1996-10-10
58	rusty	10	35	22	101	1996-10-10
22	dustin	7	45	58	103	1996-11-12
31	lubber	8	55.5	58	103	1996-11-12
58	rusty	10	35	58	103	1996-11-12

Joins

- ❖ Equi-join: special case where condition c contains only equality conjuncts (“and”)
 - ❖ Expressed as $S1 \bowtie_{sid} R1$. Implies: $S1.sid = R1.sid$
- ❖ Result schema: same as join but only one copy of sid
- ❖ Natural Join: Equi-join on all *common fields*

$S1 \bowtie_{sid} R1$

sid	sname	rating	age	bid	day
22	dustin	7	45	101	1996-10-10
58	rusty	10	35	103	1996-11-12

“Division”

- ❖ Find sailors who have reserved all boats
- ❖ Let A have 2 fields, x and y ; B have only field y :
- ❖ $A / B = \{x \mid \exists x, y \in A \forall y \in B\}$
 - ❖ i.e., A / B contains all x tuples (sailors) such that for every y tuple (boat) in B , there is an xy tuple in A .
 - ❖ Or: If the set of y values (boats) associated with an x value (sailor) in A contains all y values in B , the x value is in A / B .
- ❖ In general, x and y can be any lists of fields; y is the list of fields in B , and $x \cup y$ is the list of fields of A .

Thinking about A/B

- ❖ Idea: For A/B , compute all x values that are not *disqualified* by some y value in B .
 - ❖ x value is *disqualified* if by attaching y value from B , we obtain an xy tuple that is not in A
 - ❖ for an $x \in \pi_x(A)$ to qualify we need $\{x\} \times B$ to be in A
 - ❖ So if any $\{x\} \times B$ is not in A then that x value gets disqualified
 - ❖ Thus all disqualified x values: $\pi_x(\pi_x(A) \times B - A)$
 - ❖ $A/B : \pi_x(A) - \pi_x(\pi_x(A) \times B - A)$

Why “Division”?

- ❖ Integers:

- ❖ $a / b = c$ means something like “ c is the largest integer such that $b * c \leq a$ ”

- ❖ Relations:

- ❖ $A / B = C$ means something like “ C is the largest relation such that $B \times C \leq A$ ”

- ❖ (here for relations, “ \leq ” means set containment.)

Summary

- ❖ The relational model has rigorously defined query languages that are simple and powerful.
- ❖ Relational algebra is more operational; useful as internal representation for query evaluation plans.
- ❖ Several ways of expressing a given query; a query optimizer should choose the most efficient version.