

SQL: Part 1

History

- ❖ Structured (English) Query Language - now SQL
 - ❖ based in Relational Calculus and Algebra
- ❖ Developed for IBM's prototype RDBMS - System R
- ❖ Early 1970's with first commercial version in 1979 or so.
- ❖ First standardized in 1986 by ANSI and then ISO
- ❖ New versions released periodically, latest 2016

Tables

sp

name	item	price
S2	P3	100
S1	P2	20
S1	P1	10
S3	P4	1000
S2	P1	11
S4	P1	9

sa

name	addr
S1	123 Any St
S3	10 State St
S2	1 Main St
S4	11 State St

Looking at a table

- ❖ Table (relation) *sp*
- ❖ The * implies all columns of the table

```
select * from sp;
name | item | price
-----+-----+-----
S2   | P3   | 100
S1   | P2   | 20
S1   | P1   | 10
S3   | P4   | 1000
S2   | P1   | 11
S4   | P1   | 9
(6 rows)
```

“Selecting” a few columns

- ❖ Get name, item from table *sp*
- ❖ $\pi_{\text{name, item}}(\text{sp})$

```
select name, item
from sp;
```

name		item
-----+-----		
S2		P3
S1		P2
S1		P1
S3		P4
S2		P1
S4		P1

(6 rows)

Projection vs. SQL Select

	select name from sp;
	name

❖ Get name from table <i>sp</i>	S2
	S1
❖ $\pi_{\text{name}}(\text{sp})$ — not quite	S1
	S3
❖ No duplicate elimination	S2
unless explicitly (or implicitly)	S4
specified	(6 rows)

Projection vs. SQL Select

- ❖ Get (unique) name from table *sp*
 - ❖ $\pi_{\text{name}}(\text{sp})$
 - ❖ No duplicate elimination unless explicitly (or implicitly) specified
 - ❖ No ordering either
- ```
select distinct name
from sp;
 name

 S1
 S2
 S4
 S3
(4 rows)
```



# Relational $\sigma$ and WHERE

- ❖ Get all rows (tuples) for item P1
- ❖ The \* implies all columns of the table
- ❖  $\sigma_{\text{item}='P1'}(\text{sp})$

```
select *
from sp
where item = 'P1';
name | item | price
-----+-----+-----
S1 | P1 | 10
S2 | P1 | 11
S4 | P1 | 9
(3 rows)
```



---

# What kind of conditions

---

- ❖ comparison:  $=$ ,  $<$ ,  $>$ ,  $<>$ ,  $>=$ ,  $<=$ 
  - ❖ C1 between v1 and v2 is same as saying  
 $C1 >= v1$  and  $C1 <= v2$
- ❖ string comparison allows “LIKE” with limited pattern matching
- ❖ boolean: AND, OR, NOT
- ❖ inclusion, exclusion: IN, NOT IN



# WHERE clause example

- ❖ Get all rows (tuples) for items P1 and P2 where supplier name ends in 1.

```
select *
from sp
where item in ('P1','P2') and
 name like '%1';
```

| name | item | price |
|------|------|-------|
| S1   | P2   | 20    |
| S1   | P1   | 10    |

(2 rows)



---

# JOIN

---

- ❖ Tables SP, and, SA were derived from decomposition of the Supplier table: S(name, address, item, price)
- ❖ So to get back the original S, we need to join SP and SA on the common column “name”

```
select *
from sa join sp on sa.name = sp.name;
```

| <i>name</i> | <i>addr</i> | <i>name</i> | <i>item</i> | <i>price</i> |
|-------------|-------------|-------------|-------------|--------------|
| S2          | 1 Main St   | S2          | P3          | 100          |
| S1          | 123 Any St  | S1          | P2          | 20           |
| S1          | 123 Any St  | S1          | P1          | 10           |
| S3          | 10 State St | S3          | P4          | 1000         |
| S2          | 1 Main St   | S2          | P1          | 11           |
| S4          | 11 State St | S4          | P1          | 9            |

(6 rows)



# JOIN - alternative syntax

- ❖ A common way to specify a join is via a condition in the WHERE clause connecting the two (or more) tables
- ❖ can refer to columns of just one table using <tablename>.\*
- ❖ column names when referenced must be uniquely identifiable

```
select sa.*, item, price
from sa, sp
where sa.name = sp.name;
```

| name | addr        | item | price |
|------|-------------|------|-------|
| S2   | 1 Main St   | P3   | 100   |
| S1   | 123 Any St  | P2   | 20    |
| S1   | 123 Any St  | P1   | 10    |
| S3   | 10 State St | P4   | 1000  |
| S2   | 1 Main St   | P1   | 11    |
| S4   | 11 State St | P1   | 9     |

(6 rows)



---

# Example: JOIN + WHERE

---

- ❖ Get name and address of suppliers of item P1

```
select sa.*
from sa, sp
where sa.name = sp.name and
 item = 'P1';
```

| name        |  | addr        |
|-------------|--|-------------|
| -----+----- |  |             |
| S1          |  | 123 Any St  |
| S2          |  | 1 Main St   |
| S4          |  | 11 State St |

(3 rows)



# Expressions

```
select sa.name || ', ' || addr as Supplier,
 item, price * 0.9 as DiscountPrice
from sa, sp
where sa.name = sp.name and price * 0.9 > 10;
```

| supplier        | item | discountprice |
|-----------------|------|---------------|
| S1, 123 Any St  | P2   | 18.0          |
| S3, 10 State St | P4   | 900.0         |
| S2, 1 Main St   | P3   | 90.0          |

(3 rows)

- ❖ Pretty much any (scalar) function/operator that's applicable to the type of data is available in SQL
- ❖ If not, you can create your own function
- ❖ WHERE conditions can also have expressions
- ❖ "as": alias ("rename operator") to give a new name to a column or expression



# Subqueries

- ❖ Find suppliers who supplied part P1
- ❖ The IN operator checks if the value(s) are in the *subquery*
- ❖ Can be powerful, if you can keep track of the meaning
- ❖ If the subquery has no references to “outside”, then it’s called *uncorrelated*

```
select *
from sa
where name in
 (select name
 from sp
 where item = 'P1');
```

| name | addr        |
|------|-------------|
| S1   | 123 Any St  |
| S2   | 1 Main St   |
| S4   | 11 State St |

(3 rows)



---

# Subqueries

---

- ❖ Find suppliers who supplied part P1
- ❖ Another way to express the same request using EXISTS and a *correlated* subquery
- ❖ EXISTS: the relation specified is a non-empty relation
- ❖ Many possibilities: use intuition and keep it simple

```
select *
from sa
where exists
 (select name
 from sp
 where sp.name = sa.name
 and item = 'P1');
```

| name | addr        |
|------|-------------|
| S1   | 123 Any St  |
| S2   | 1 Main St   |
| S4   | 11 State St |

(3 rows)



---

# Simple Alternative?

---

- ❖ Subqueries can on many occasions be expressed as simple JOIN query

```
select sa.*
from sa, sp
where sa.name = sp.name
 and item = 'P1';
```

| name  |   | addr        |
|-------|---|-------------|
| ----- | + | -----       |
| S1    |   | 123 Any St  |
| S2    |   | 1 Main St   |
| S4    |   | 11 State St |

(3 rows)



---

# Subqueries

---

- ❖ [IN | NOT IN] <subq> - set membership check
- ❖ [EXISTS | NOT EXISTS | UNIQUE | NOT UNIQUE] <subq>
- ❖ <value> *relop* [ANY | ALL] <subq>
  - ❖ *relops* are >, <, =, etc.
  - ❖ <subq> must return 1 row unless ANY / ALL used



# Cross Product

- ❖ Given a good relational model, need for cross products should be rare
- ❖ Mostly for understanding SQL processing
- ❖ adding a proper WHERE condition makes a cross product into a join

```
select *
from sa, sp;
```

| name | addr        | name | item | price |
|------|-------------|------|------|-------|
| S1   | 123 Any St  | S2   | P3   | 100   |
| S3   | 10 State St | S2   | P3   | 100   |
| S2   | 1 Main St   | S2   | P3   | 100   |
| S4   | 11 State St | S2   | P3   | 100   |
| S1   | 123 Any St  | S1   | P2   | 20    |
| S3   | 10 State St | S1   | P2   | 20    |
| S2   | 1 Main St   | S1   | P2   | 20    |
| S4   | 11 State St | S1   | P2   | 20    |
| S1   | 123 Any St  | S1   | P1   | 10    |
| S3   | 10 State St | S1   | P1   | 10    |
| S2   | 1 Main St   | S1   | P1   | 10    |
| S4   | 11 State St | S1   | P1   | 10    |
| S1   | 123 Any St  | S3   | P4   | 1000  |
| S3   | 10 State St | S3   | P4   | 1000  |
| S2   | 1 Main St   | S3   | P4   | 1000  |
| S4   | 11 State St | S3   | P4   | 1000  |
| S1   | 123 Any St  | S2   | P1   | 11    |
| S3   | 10 State St | S2   | P1   | 11    |
| S2   | 1 Main St   | S2   | P1   | 11    |
| S4   | 11 State St | S2   | P1   | 11    |
| S1   | 123 Any St  | S4   | P1   | 9     |
| S3   | 10 State St | S4   | P1   | 9     |
| S2   | 1 Main St   | S4   | P1   | 9     |
| S4   | 11 State St | S4   | P1   | 9     |

(24 rows)



---

# Review

---

SELECT [options] column\_expression\_list or "\*"
FROM table\_expression\_list
WHERE condition

- ❖ We are SELECTing a set of expressions forming columns of the output
  - ❖ RA calls it projection
- ❖ The rows that contribute to it come from the tables, or equivalent, in the FROM clause
  - ❖ RA cross product (aka cross join)
- ❖ But only the rows that meet the WHERE condition
  - ❖ RA calls it selection
  - ❖ Some conditions may be specifying the RA natural join



---

# Conceptually

---

- ❖ Compute the cross product of relations in FROM
- ❖ Keep the rows that meet the WHERE conditions
- ❖ WHERE conditions comprise
  - ❖ join conditions (converting cross product to joins)
  - ❖ conditions on attributes, including subqueries
  - ❖ (subqueries are a kind of join...)
- ❖ Compute the expressions using values from rows (and literals and maths and scalar function and ...)



---

# DISTINCT option

---

```
SELECT [options] column_expression_list or "*"
FROM table_expression_list
WHERE condition
```

- ❖ One of the options is “**DISTINCT**”
- ❖ Applied after all the output rows are formed
- ❖ Removes duplicate rows in the output (and converts it into a true *relation*)
- ❖ Not used much
  - ❖ information loss
  - ❖ shouldn't need it much in a well-defined schema
  - ❖ aggregation queries don't need it



---

# DISTINCT option

---

- ❖ Get all the distinct items being supplied

```
select distinct item
from sp;
```

```
 item

 P2
 P1
 P4
 P3
(4 rows)
```



---

# SET operations

---

- ❖ UNION [ALL]
- ❖ INTERSECT [ALL]
- ❖ EXCEPT [ALL] (set minus)
- ❖ The table\_expressions have to be compatible: i.e. must have same number of columns and each corresponding column must have compatible type
- ❖ without ALL, no duplicates (even if original table\_expressions have them)
- ❖ with ALL
  - ❖ all rows kept in UNION
  - ❖ matching number of duplicates kept in INTERSECT
  - ❖ matching number of duplicates discarded in EXCEPT



---

# INTERSECT

---

- ❖ Suppliers supplying item P1 and P2
  - ❖ Suppliers for P1
  - ❖ intersect
  - ❖ Suppliers for P2

```
select sa.*
from sa, sp
where sa.name = sp.name
 and item = 'P1'
intersect
select sa.*
from sa, sp
where sa.name = sp.name
 and item = 'P2';
```

| name | addr       |
|------|------------|
| S1   | 123 Any St |

(1 row)



# EXCEPT – set minus

- ❖ Suppliers supplying item P1 but not P2
- ❖ Suppliers for P1
- ❖ except
- ❖ Suppliers for P2

```
select sa.*
from sa, sp
where sa.name = sp.name
 and item = 'P1'
```

```
except
select sa.*
from sa, sp
where sa.name = sp.name
 and item = 'P2';
```

| name | addr        |
|------|-------------|
| S2   | 1 Main St   |
| S4   | 11 State St |

(2 rows)



---

# UNION

---

❖ Suppliers supplying item P1 or P2

❖ Suppliers for P1

❖ union

❖ Suppliers for P2

❖ Suppliers for either item i.e.

❖ item IN (P1, P2) or

❖ item = P1 or item = P2.

```
select sa.*
from sa, sp
where sa.name = sp.name
 and item = 'P1'
```

```
union
select sa.*
from sa, sp
where sa.name = sp.name
 and item = 'P2';
```

| name | addr        |
|------|-------------|
| S2   | 1 Main St   |
| S1   | 123 Any St  |
| S4   | 11 State St |

(3 rows)



---

# UNION

---

❖ Suppliers supplying item P1 or P2

❖ Suppliers for P1

❖ union

❖ Suppliers for P2

❖ Suppliers for either item i.e.

❖ item IN (P1, P2) or

❖ item = P1 or item = P2.

```
select distinct sa.*
from sa, sp
where sa.name = sp.name
 and item in ('P1', 'P2')
```

```
;
```

| name | addr        |
|------|-------------|
| S2   | 1 Main St   |
| S1   | 123 Any St  |
| S4   | 11 State St |

(3 rows)



---

# ORDER BY

---

- ❖ Notice that almost none of the results in our examples are “ordered”
- ❖ This stems from the SET foundations of RDBMS
- ❖ ORDER BY clause helps us order results in the output

```
select sa.*, item, price
from sa, sp
where sa.name = sp.name;
```

| name | addr        | item | price |
|------|-------------|------|-------|
| S2   | 1 Main St   | P3   | 100   |
| S1   | 123 Any St  | P2   | 20    |
| S1   | 123 Any St  | P1   | 10    |
| S3   | 10 State St | P4   | 1000  |
| S2   | 1 Main St   | P1   | 11    |
| S4   | 11 State St | P1   | 9     |

(6 rows)



# ORDER BY

```
select sa.*, item, price
from sa, sp
where sa.name = sp.name
order by sa.name, item;
```

- ❖ ORDER BY clause helps us order results in the output
- ❖ Options are ASC (default), DESC, NULLS FIRST, NULLS LAST

| name | addr        | item | price |
|------|-------------|------|-------|
| S1   | 123 Any St  | P1   | 10    |
| S1   | 123 Any St  | P2   | 20    |
| S2   | 1 Main St   | P1   | 11    |
| S2   | 1 Main St   | P3   | 100   |
| S3   | 10 State St | P4   | 1000  |
| S4   | 11 State St | P1   | 9     |

(6 rows)