

Relational Model

Linda Wu

(CMPT 354 • 2004-2)

Topics

- Relational model
- SQL language
- Integrity constraints
- ER to relational
- Views

Why Study the Relational Model?

- Most widely used model
 - Vendors: IBM, Informix, Microsoft, Oracle, Sybase, etc.
- Recent competitor: object-oriented model
 - ObjectStore, Versant, Ontos
- A synthesis emerging: object-relational model
 - Informix Universal Server, UniSQL, O2, Oracle, DB2

Relational Model

- Relational database: a set of relations
- A relation is made up of 2 parts
 - Instance: a table, with rows and columns
 - # of Rows = cardinality
 - # of fields = degree / arity
 - Schema: specifies name of relation, plus name and domain of each column
e.g., Students (sid: string, name: string, login: string, age: integer, gpa: real)
- A relation can be thought of as a set of **unique** rows or tuples

Relational Model (Cont.)

<i>sid</i>	<i>name</i>	<i>login</i>	<i>age</i>	<i>gpa</i>
S0001	Jones	jones@cs	18	3.4
S0002	Smith	smith@ee	19	3.2
S0003	Smith	smith@math	18	3.8
S0004	Mary	mary@music	14	1.8

Cardinality = 4, degree = 5, all rows distinct

Relational Model (Cont.)

- A major strength of relational model
 - Supports simple, powerful querying of data
- Queries can be written intuitively, and the DBMS is responsible for efficient evaluation
 - The key: precise semantics for relational queries
 - Allows the optimizer to extensively re-order operations, and still ensure that the answer does not change

SQL Language

- Query
 - A question about the data
 - Its answer consists of a relation containing the result
- Query language
 - A specialized language for writing query
- SQL (Structured Query Language)
 - The standard language for creating, manipulating, and querying data in a relational DBMS
 - Originally developed by IBM (system-R) in the 1970s

SQL Language (Cont.)

- Need for a standard since SQL products are offered by many vendors
 - SQL-86
 - SQL-89 (minor revision)
 - SQL-92 (major revision)
 - SQL-99 (major extension, current standard)

SQL Language (Cont.)

Querying single relation

E.g., to find all 18 year old students:

```
SELECT *  
FROM Students S  
WHERE S.age=18
```

```
SELECT S.name, S.login  
FROM Students S  
WHERE S.age=18
```

Instance of Students

sid	name	login	age	gpa
S0001	Jones	jones@cs	18	3.4
S0002	Smith	smith@ee	19	3.2
S0003	Smith	smith@math	18	3.8
S0004	Mary	mary@music	14	1.8
S0005	Gary	gary@biz	12	2.0

Chapter 3

9

SQL Language (Cont.)

Querying multiple relations

```
SELECT S.name, E.cid  
FROM Students S, Enrolled E  
WHERE S.sid=E.sid AND E.grade='A'
```

Instance of Enrolled

sid	cid	grade
S0002	CMPT101	C
S0003	ECON205	B
S0004	BUS310	A
S0005	CMPT250	B

The query result is:

S.name	E.cid
Mary	BUS310

Chapter 3

CMPT 354 • 2004-2

10

SQL Language (Cont.)

Creating relations

- SQL uses the word *table* to denote relation
- The subset of SQL that supports the creation, deletion and modification of tables is called Data Definition Language (DDL)
- The type (domain) of each field is specified, and enforced by the DBMS whenever tuples are added or modified

```
CREATE TABLE Enrolled  
(sid CHAR(20),  
cid CHAR(20),  
grade CHAR(2))
```

```
CREATE TABLE Students  
(sid CHAR(20),  
name CHAR(20),  
login CHAR(10),  
age INTEGER,  
gpa REAL)
```

Chapter 3

CMPT 354 • 2004-2

11

SQL Language (Cont.)

Adding and deleting Tuples

Insert tuples

```
INSERT INTO Students (sid, name, login, age, gpa)  
VALUES (S0002, 'Smith', 'smith@ee', 19, 3.2)  
OR  
INSERT INTO Students  
VALUES (S0002, 'Smith', 'smith@ee', 19, 3.2)
```

Delete all tuples satisfying some conditions

```
DELETE  
FROM Students S  
WHERE S.name = 'Smith'
```

Chapter 3

CMPT 354 • 2004-2

12

SQL Language (Cont.)

○ Modifying tuple values

Modify the column values in an existing row

```
UPDATE Students S
SET S.age = S.age + 1, S.gpa = S.gpa - 1
WHERE S.sid = 'S0003'
```

- WHERE clause is applied first and determine the rows to modify
- SET clause then determines how to modify the rows; the column value at the right side of “=” is the old value

SQL Language (Cont.)

○ Destroying and Altering Relations

```
DROP TABLE Students RESTRICT/CASCADE
```

Destroys the relation Students: the schema information *and* the tuples are deleted

```
ALTER TABLE Students
ADD COLUMN firstYear INTEGER
```

The schema of Students is altered by adding a new field; every tuple in the current instance is extended with a *null* value in the new field

Integrity Constraints

○ Integrity constraint (IC)

- Condition that must be true for any instance of the database
- Specified when schema is defined
- Checked when relations are modified
- E.g., domain constraints, primary/foreign key constraints

○ A legal instance of a relation is one that satisfies all specified ICs

- DBMS should not allow illegal instances

Primary Key Constraints

○ A set of fields is a (candidate) key for a relation if:

1. No two distinct tuples have same values in all key fields, and,
2. No subset of the key fields is a unique identifier for a tuple

- Condition 2 false? A superkey
- If there's more than one key for a relation, one of the keys is chosen (by DBA) to be the **primary key**

○ Example: in Students relation

- *sid* is a key; *name* is not a key
- The set {*sid*, *gpa*} is a superkey

Primary Key Constraints (Cont.)

- Specifying key constraints in SQL
 - candidate keys specified using UNIQUE
 - primary key specified using PRIMARY KEY

```
CREATE TABLE Enrolled
( sid   CHAR(20),
  cid   CHAR(20),
  grade CHAR(2),
  PRIMARY KEY (sid, cid) )
```

For a given student and course,
there is a single grade.

```
CREATE TABLE Enrolled
( sid   CHAR(20),
  cid   CHAR(20),
  grade CHAR(2),
  PRIMARY KEY (sid),
  UNIQUE (cid, grade) )
```

Students can take only one
course, and receive a single grade
for that course; no two students in
a course receive the same grade.

Primary Key Constraints (Cont.)

- Enforcing primary key constraints
 - Insertion and update can cause violations
 - Deletion does not cause a violation
 - Examples
 - Insert (**S0002**, Mike, mike@cs, 20, 2.7) to Students
 - Insert (**null**, Mike, mike@cs, 20, 2.7) to Students
 - Change sid "S0001" to "**S0002**" in the first tuple of Students

Foreign Key Constraints

- Foreign key
 - A set of fields in one relation that is used to "refer" to a tuple in another relation like a "logical pointer"
 - The foreign key in the referencing relation must match the primary key of the referenced relation

E.g., Enrolled (*sid*: string, *cid*: string,
grade: string)

- *sid* is a foreign key referring to Students
- If all foreign key constraints are enforced, referential integrity is achieved, i.e., no dangling references

Foreign Key Constraints (Cont.)

- Foreign Keys in SQL

```
CREATE TABLE Enrolled
( sid   CHAR(20),
  cid   CHAR(20),
  grade CHAR(2),
  PRIMARY KEY (sid, cid),
  FOREIGN KEY (sid) REFERENCES Students )
```

- Only students listed in the Students relation are allowed to enroll in courses
- A student has exactly one grade for each course s/he is enrolled in

Foreign Key Constraints (Cont.)

Enforcing Referential Integrity

Foreign key			Primary key				
<i>cid</i>	<i>grade</i>	<i>sid</i>	<i>sid</i>	<i>name</i>	<i>login</i>	<i>age</i>	<i>gpa</i>
CMPT101	C	S0002	S0001	Jones	jones@cs	18	3.4
ECON205	B	S0003	S0002	Smith	smith@ee	19	3.2
BUS310	A	S0004	S0003	Smith	smith@math	18	3.8
CMPT250	B	S0005	S0004	Mary	mary@music	14	1.8
			S0005	Gary	gary@biz	12	2.0

Enrolled Students

- Insert (PHYS110, A, S1234) to Enrolled (rejected!)
- Delete (S0002, Smith, smith@ee, 19, 3.2) from Students (?)
- Change sid "S005" to "S9999" in the last tuple (?)

Foreign Key Constraints (Cont.)

A foreign key could refer to the same relation

<i>sid</i>	<i>name</i>	<i>login</i>	<i>age</i>	<i>gpa</i>	<i>partner</i>
S0001	Jones	jones@cs	18	3.4	S0002
S0002	Smith	smith@ee	19	3.2	S0001
S0003	Smith	smith@math	18	3.8	null
S0004	Mary	mary@music	14	1.8	null
S0005	Gary	gary@biz	12	2.0	null

Students

a foreign key referring to Students

Foreign Key Constraints (Cont.)

- ### Referential Integrity in SQL: SQL-92/99 support 4 options on DELETE and UPDATE
- NO ACTION (default)
 - delete/update is rejected
 - CASCADE
 - Also delete/update all tuples that refer to deleted/updated tuple
 - SET NULL
 - Set foreign key value of the referencing tuple to null
 - SET DEFAULT
 - Set foreign key value of the referencing tuple to a default value

Foreign Key Constraints (Cont.)

```
CREATE TABLE Enrolled
( sid    CHAR(20) DEFAULT 'S9999',
  cid    CHAR(20),
  grade  CHAR(2),
  PRIMARY KEY (sid, cid),
  FOREIGN KEY (sid) REFERENCES Students
    ON DELETE CASCADE
    ON UPDATE SET DEFAULT )
```

Where do ICs Come From?


- ICs are based upon the semantics of the real-world enterprise that is being described in the database relations
- We can check a database instance to see if an IC is violated, but we can **NEVER** infer that an IC is true by looking at an instance
 - An IC is a statement about *all possible* instances!
- Key and foreign key ICs are the most common
- More general ICs are supported too

Logical DB Design: ER to Relational

- Entity set
- Relationship set without constraints
- Relationship set with key constraints
- Relationship set with participation constraints
- Weak entity set
- Class hierarchy
- Aggregation

Entity sets to Tables

- Attributes of entity set → attributes of table
- Domain of each attribute
- Primary key of entity set




```
CREATE TABLE Employees
( ssn CHAR(11),
  name CHAR(20),
  lot INTEGER,
  PRIMARY KEY (ssn) )
```

ssn	name	lot
123-22-3666	Andrew	44
231-31-5368	Emily	22
131-24-3650	Smith	15

Relationship Sets to Tables

- In translating a relationship set without key or participation constraints to a relation, attributes of the relation must include:
 - The primary key attributes for each participating entity set, as foreign keys
 - This set of attributes forms a superkey for the relation
 - If no key constraints, this set of attributes is a candidate key
 - All descriptive attributes

Relationship Sets to Tables (Cont.)




```
CREATE TABLE Works_In
(
  ssn CHAR(1),
  did INTEGER,
  since DATE,
  PRIMARY KEY (ssn, did),
  FOREIGN KEY (ssn) REFERENCES Employees,
  FOREIGN KEY (did) REFERENCES Departments )
```

Chapter 3

CMPT 354 • 2004-2

29

Relationship Sets to Tables (Cont.)



```
CREATE TABLE Reports_To
(
  supervisor_ssn CHAR(11),
  subordinate_ssn CHAR(11),
  PRIMARY KEY (supervisor_ssn, subordinate_ssn),
  FOREIGN KEY (supervisor_ssn) REFERENCES Employees (ssn),
  FOREIGN KEY (subordinate_ssn) REFERENCES Employees (ssn) )
```


Chapter 3

CMPT 354 • 2004-2

30

Translating Relationship Set with Key Constraints

- Map Manages relationship to a table
 - did* is the key now
 - Separate tables for Employees and Departments



```
CREATE TABLE Manages (
  ssn CHAR(11),
  did INTEGER,
  since DATE,
  PRIMARY KEY (did),
  FOREIGN KEY (ssn) REFERENCES Employees,
  FOREIGN KEY (did) REFERENCES Departments )
```

Chapter 3

Translating Relationship Set with Key Constraints (Cont.)

- Since each department has a unique manager, we could instead combine Manages and Departments
 - Avoid a distinct table for the relationship set, therefore, fast query
 - Space would be wasted if some departments have no managers (null)

```
CREATE TABLE Dept_Mgr (
  did INTEGER,
  dname CHAR(20),
  budget REAL,
  ssn CHAR(11),
  since DATE,
  PRIMARY KEY (did),
  FOREIGN KEY (ssn) REFERENCES Employees)
```

Chapter 3

32

Translating Relationship Set with Participation Constraints

- Every department has a manager
 - Use a combined table for Manages and Departments

```
CREATE TABLE Dept_Mgr (
    did    INTEGER,
    dname  CHAR(20),
    budget REAL,
    ssn    CHAR(11) NOT NULL,
    since  DATE,
    PRIMARY KEY (did),
    FOREIGN KEY (ssn) REFERENCES Employees
    ON DELETE NO ACTION)
```

- If wish to delete an Employees tuple: first change the Dept_Mgr tuple to have a new employee as manager

Translating Relationship Set with Participation Constraints (Cont.)


- If use a separate table for Manages:
 - Does not ensure that a manager is initially appointed for each department

```
CREATE TABLE Manages (
    ssn    CHAR(11) NOT NULL,
    did    INTEGER,
    since  DATE,
    PRIMARY KEY (did),
    FOREIGN KEY (ssn) REFERENCES Employees,
    FOREIGN KEY (did) REFERENCES Departments)
```

Combined table (Dept_Mgr) is better than separate table (Manages) for one-to-many relationship, especially when the entity set with key constraint also has a total participation constraint

Translating Relationship Set with Participation Constraints (Cont.)

- We can capture participation constraints involving one entity set in a binary relationship, but little else (without resorting to CHECK constraints)




Ensuring total participation of Departments in Works_In

- Possible way: declares that *did* in Departments is a foreign key referring to Works_In
- Problem: this is not a valid FK constraint, since *did* is not a key in Works_In
- Solution: assertion

Translating Weak Entity Set

- A weak entity can be identified uniquely only by considering the primary key of another (owner) entity
 - Owner entity set and weak entity set participate in a one-to-many relationship set: 1 owner, many weak entities
 - Weak entity set have total participation in its identifying relationship set
- Weak entity set and identifying relationship set are translated into a single table
 - When the owner entity is deleted, all owned weak entities must also be deleted

Translating Weak Entity Set (Cont.)



```
CREATE TABLE Dep_Policy (
  pname CHAR(20),
  age INTEGER,
  cost REAL,
  ssn CHAR(11) NOT NULL,
  PRIMARY KEY (pname, ssn),
  FOREIGN KEY (ssn) REFERENCES Employees,
  ON DELETE CASCADE)
```

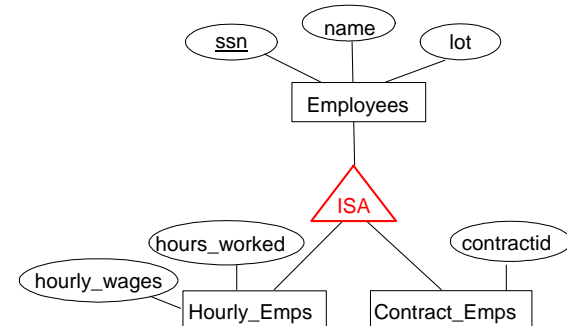
Chapter 3

CMPT 354 • 2004-2

37

Translating Class Hierarchies

- Example: Hourly_Emps and Contract_Emps



Chapter 3

CMPT 354 • 2004-2

38

Translating Class Hierarchies (Cont.)

- General approach: 3 distinct relations
 - Employees (ssn, name, lot)
 - All employees are recorded
 - Hourly_Emps (ssn, hourly_wages, hours_worked)
 - ssn is FK referencing Employees
 - Must delete Hourly_Emps tuple if the referenced Employees tuple is deleted
 - Contract_Emps (ssn, contractid)
- * Queries involving all employees easy
- * Queries involving just Hourly_Emps require a join with Employees to get some attributes

Chapter 3

CMPT 354 • 2004-2

39

Translating Class Hierarchies (Cont.)

- Alternative: 2 relations
 - Hourly_Emps: (ssn, name, lot, hourly_wages, hours_worked)
 - Contract_Emps: (ssn, name, lot, contractid)
- * Each employee must be in one of these two subclasses
- * Duplication: name and lot are stored twice if an employee is both an Hourly_Emps and a Contract_Emps
- * Queries involving all employees have to examine two relations

Chapter 3

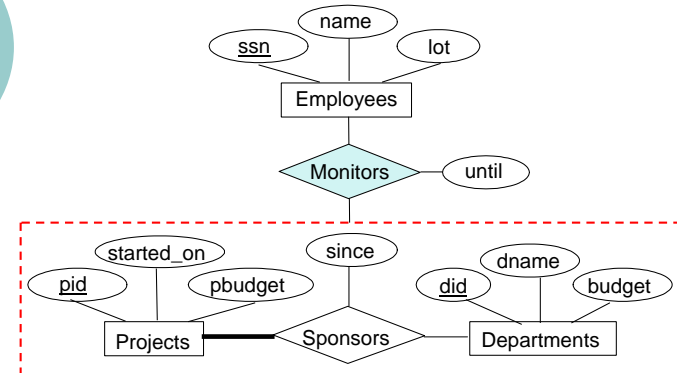
CMPT 354 • 2004-2

40

Translating Aggregation


- Use the standard mapping for a relationship set
 - The primary key attributes of each participating entity set
 - The descriptive attributes of the relationship set
- Example: Monitors
 - 5 relations: Employees, Projects, Departments, Sponsors, Monitors (*ssn*, *did*, *pid*, *until*)
 - Sponsors can be dropped if it:
 - Has no descriptive attributes, and,
 - Has total participation in Monitors

Translating Aggregation (Cont.)



Translating Ternary Relationship

- Key constraints
 - Combine Purchaser with Policies
 - Combine Beneficiary with Dependents
- Participation constraints
 - Lead to NOT NULL constraints
- Weak entity set
 - Combine Dependents and Beneficiary



Translating Ternary Relationship (Cont.)

```
CREATE TABLE Policies (
    policyid INTEGER,
    cost REAL,
    ssn CHAR(11) NOT NULL,
    PRIMARY KEY (policyid),
    FOREIGN KEY (ssn) REFERENCES Employees,
    ON DELETE CASCADE )
```

```
CREATE TABLE Dependents (
    pname CHAR(20),
    age INTEGER,
    policyid INTEGER,
    PRIMARY KEY (pname, policyid),
    FOREIGN KEY (policyid) REFERENCES Policies,
    ON DELETE CASCADE )
```

Views

- A view is just a relation, but we store a definition of it, rather than a set of tuples

```
CREATE VIEW YoungActiveStudents (name, grade)
AS SELECT S.name, E.grade
FROM Students S, Enrolled E
WHERE S.sid = E.sid and S.age < 21
```

- * To drop a view:
DROP VIEW command
- * To drop a table when there is a view on the table:
DROP TABLE command with **CASCADE** option

Views (Cont.)

- View provides support for logical data independence
- Views provides support for security
 - View can be used to present necessary information (or a summary), while hiding details in underlying relations
- We can insert a row into a view by inserting a row into its base table
- An INSERT or UPDATE to the base table may lead to a row not in the view

Relational Model: Summary

- A tabular representation of data
- Simple and intuitive, the most widely used
- Integrity constraints can be specified by the DBA, based on application semantics; DBMS checks for violations
 - Two important ICs: primary and foreign keys
 - In addition, we always have domain constraints
- Powerful and natural query languages exist
- Rules to translate ER to relational model