

Hashed-Based Indexing

Linda Wu

(CMPT 354 • 2004-2)


Topics

- Static hashing
 - Extendible Hashing
 - Linear Hashing
- Dynamic hashing

Static Hashing

- An index consists of buckets $0 \sim N-1$
- A bucket consists of one primary page and, possibly, additional overflow pages
- Buckets contain data entries
- Hash function h
 - h must distribute values in the domain of search field uniformly over the buckets
 - $h(k) = (a * k + b)$ usually works well, a and b are constants for tuning h
 - $h(k) \bmod N$: bucket to which the data entry with search key k belongs

Static Hashing (Cont.)



- # of primary pages, N , is fixed
- Primary pages are allocated sequentially, never de-allocated; overflow pages allocated if needed
- Long overflow chains can develop and degrade performance

Static Hashing (Cont.)

- Search
 - Identify the correct bucket using h
 - Search the bucket for the data entry
- Insert
 - Identify the correct bucket using h
 - If no space in the bucket, allocate a new overflow page to the overflow chain of the bucket, put the data entry on this page
- Delete
 - Identify the correct bucket using h
 - Search the bucket for the data entry, remove it
 - If it is the last in an overflow page, remove the page from the chain and de-allocate the page

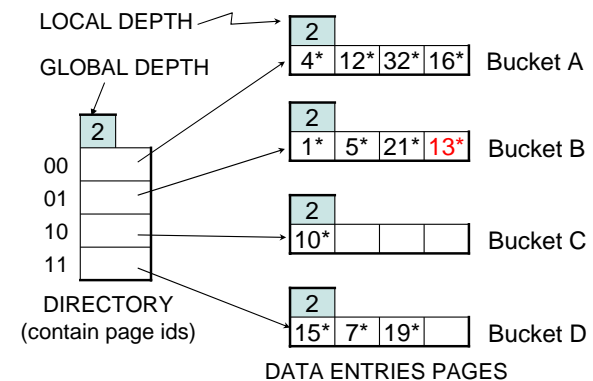
Extendible Hashing

- To avoid overflow page
 - Possible solution: if a bucket is full, reorganize file by doubling # of buckets and redistributing the entries across the new set of buckets
 - * *Reading and writing all pages is expensive!*
- Solution
 - Use a directory of pointers to buckets, double # of buckets by doubling the directory, splitting just the bucket that overflowed!
 - * Directory is much smaller than file, so doubling it is much cheaper. Only one page of data entries is split. No overflow page!
 - * Trick lies in how hash function is adjusted!

Extendible Hashing (Cont.)

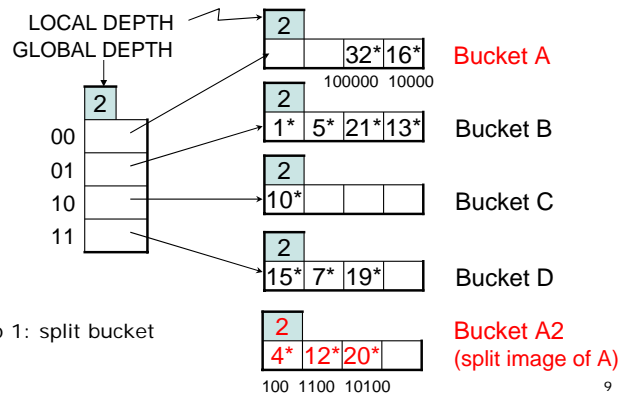
- Example
 - Directory is an array of size 4
 - To find the correct bucket for k
 - Take last **global depth** (2) bits of $h(k)$
 - e.g., if $h(k) = 5 = 101$ (binary), it is in the bucket pointed to by dictionary element 01
 - To insert
 - If bucket is full, split it (allocate a new page, re-distribute entries)
 - If necessary, double the directory

Extendible Hashing (Cont.)



Extendible Hashing (Cont.)


- Insert entry r with $h(r) = 20$



Chapter 11

9

Extendible Hashing (Cont.)



Chapter 11

10

Extendible Hashing (Cont.)

- Points to note
 - Allocate a new bucket page; write both this page and the old bucket page; double the dictionary array
 - $h(r) = 20 = \text{binary } 10100$: last 2 bits (00) tell r belongs in A or A2; last 3 bits (000 / 1000) are needed to tell which bucket
- Global depth of directory
 - Max # of bits needed to tell which bucket an entry belongs to
- Local depth of a bucket
 - # of bits used to determine if an entry belongs to this bucket

Chapter 11

CMPT 354 • 2004-2

11

Extendible Hashing (Cont.)

- When does bucket split cause directory doubling?
 - Before inserting, local depth of bucket = global depth
 - Inserting causes local depth to become > global depth
- Directory is doubled by copying it over and adjusting pointer to split image page (use of least significant bits enables efficient doubling via copying of directory!)

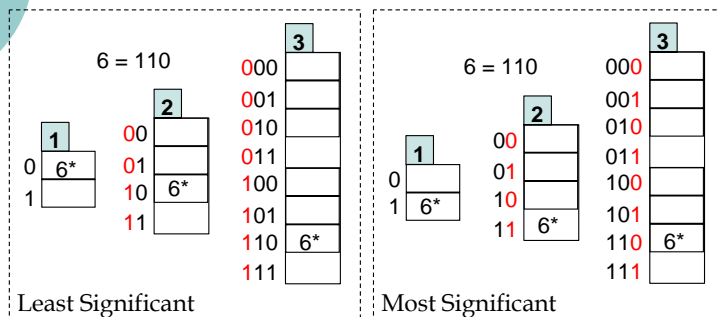
Chapter 11

CMPT 354 • 2004-2

12

Extendible Hashing (Cont.)

- Why use least significant bits in directory?
 - Allows for efficient doubling via copying



Extendible Hashing (Cont.)

- Delete
 - Locate the data entry by computing its hash value, taking the last bits, and looking in the bucket pointed to by this directory element
 - Remove the data entry
 - If the removal makes the bucket empty, it can be merged with its split image; local depth is decreased
 - If each directory element points to same bucket as its split image, the directory can be halved; global depth is decreased

* The last 2 steps can be omitted

Extendible Hashing (Cont.)

- I/O cost of equality search
 - If the directory fits in memory, equality search can be answered with one disk access
 - otherwise, two
- Collision (duplicate) handling
 - Collision: multiple data entries with the same hash value
 - Use overflow page when more data entries than will fit on a page have the same hash value

Linear Hashing


- An alternative to Extendible Hashing
- Not require directory; handle duplicates
- Idea: use a family of hash functions h_0, h_1, h_2, \dots
 - $h_i(\text{key}) = h(\text{key}) \bmod (2^i N)$
 - N = initial # buckets; h is some hash function (range is not 0 to $N-1$)
 - If $N = 2^{d_0}$, h_i consists of applying h and looking at the last d_i bits, where $d_i = d_0 + i$
 - h_{i+1} doubles the range of h_i (similar to directory doubling)

Linear Hashing (Cont.)

- Linear Hashing avoids directory by using overflow pages, and choosing bucket to split round-robin
 - Splitting proceeds in "rounds"; current round number is $Level$ (initial value is 0)
 - Bucket to split is denoted by $Next$
 - $Next$ is initialized to 0 when a new round begins, and increased by 1 after a splitting
 - Buckets (0 ~ $Next-1$) have been split; buckets ($Next \sim N_{Level}$) yet to be split
 - Splitting is triggered when an overflow page is added, and $h_{Level+1}$ redistributes entries between this bucket and its split image
 - The round $Level$ ends when all N_{Level} initial buckets are split

Linear Hashing (Cont.)

- In the middle of a round



Linear Hashing (Cont.)

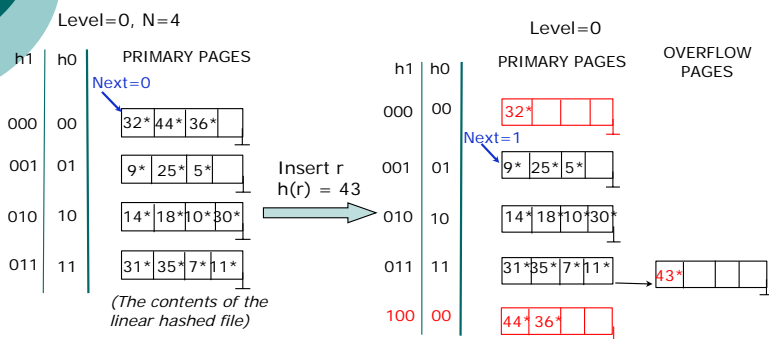
- Search for data entry r
 - To find bucket for data entry r , calculate $h_{Level}(r)$
 - If $h_{Level}(r)$ is in the range ($Next \sim N_R$), r belongs here
 - Else, r could belong to bucket $h_{Level}(r)$ or bucket $h_{Level}(r) + N_{Level}$; must apply $h_{Level+1}(r)$ to find out
- Insert
 - Find bucket by applying $h_{Level} / h_{Level+1}$
 - If bucket to insert into is full
 - Add overflow page and insert data entry
 - Split $Next$ bucket (its entries are redistributed by $h_{Level+1}$), and increment $Next$

Linear Hashing (Cont.)

- Actually any criterion can be chosen to trigger splitting
- Since buckets are split round-robin, long overflow chains don't develop!
- Doubling of directory in Extendible Hashing is similar; switching of hash functions is implicit in how the # of bits examined is increased

Linear Hashing (Cont.)

- On split, $h_{Level+1}$ is used to redistribute entries

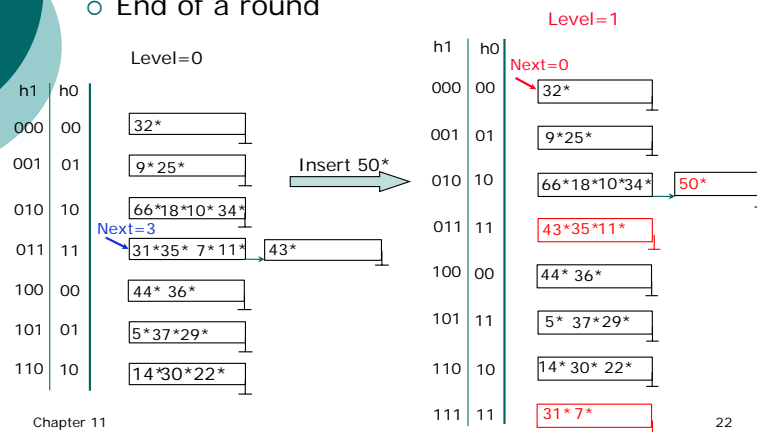


Chapter 11

21

Linear Hashing (Cont.)

- End of a round



Chapter 11

22

Extendible vs. Linear Hashing

- The two schemes are actually quite similar
 - Begin with an imaginary directory in LH with N elements
 - First split is at bucket 0 (the imaginary directory is doubled at this point). Since elements $\langle 1, N+1 \rangle$, $\langle 2, N+2 \rangle$, ... are the same, we need only create directory element N, which differs from 0, now. When bucket 1 splits, create directory element N+1, etc.
 - The directory is doubled gradually. Also, primary bucket pages are created in order. If they are allocated in sequence too, we don't need a directory!

Chapter 11

CMPT 354 • 2004-2

23

Extendible vs. Linear Hashing (Cont.)

Moving from h_i to h_{i+1} in LH corresponds to doubling the directory in EH

- Extendible Hashing
 - Directory is doubled in a single step
 - always splitting the appropriate bucket (dense data area): reduced # of splits and a higher bucket occupancy
- Linear Hashing
 - Directory is doubled gradually over the course of a round
 - A directory can be avoided by a clever choice of the buckets to split

Chapter 11

CMPT 354 • 2004-2

24