

Evaluating Relational Operators

Linda Wu

(CMPT 354 • 2004-2)

Topics

- Selection operation
- Projection operation
- Set operations
- Aggregate operations
- Join operation
- Impact of buffering

Relational Operations

- We will consider how to implement:
 1. Selection (σ): select a subset of rows from relation
 2. Projection (π): delete unwanted columns from relation
 3. Union (\cup): tuples in relation 1 or 2
 4. Set-difference ($-$): tuples in relation 1 but not in relation 2
 5. Aggregation (SUM, MIN, etc) and GROUP BY
 6. Join (\bowtie): combine two relations
 - Intersection (\cap) and cross-product (\times) are implemented as special cases of join

Schema for Examples

Sailors (*sid*: integer, *sname*: string, *rating*: integer, *age*: real)

Reserves (*sid*: integer, *bid*: integer, *day*: date, *rname*: string)

- Similar to old schema; *rname* added for variations
- Reserves
 - Each tuple is 40 bytes long; 100 tuples per page; 1000 pages
- Sailors
 - Each tuple is 50 bytes long; 80 tuples per page; 500 pages

Selection σ

$$\sigma_{R.attr \text{ op value}}(R)$$

- No index on $R.attr$, R is not sorted on $attr$
 - Scan the **entire** relation R
 - Add tuples to result if the condition is satisfied
- No index on $R.attr$, R is sorted on $attr$
 - Binary search to find the **first** qualifying tuple
 - From there, scan R till the condition is no longer satisfied
- B+ tree index on $R.attr$
 - Search the tree to find the **first** index pointing to a qualifying tuple of R
 - Scan the leaf pages to retrieve all qualifying data entries, and, the corresponding tuples if Alternative (2), (3)
- Hash index on $R.attr$, op is equality
 - Retrieve the correct bucket page in the index
 - Retrieve qualifying tuples from R

Selection σ (Cont.)

- If using an index for selections, the I/O cost depends on # of qualifying tuples and clustering

- Cost of finding qualifying data entries (typically small) + cost of retrieving tuples (could be large without clustering)
- Example

```
SELECT * FROM Reserves R
WHERE R.rname < 'C%'
```

Assuming uniform distribution of names, about 10% of tuples qualify (100 pages; 10,000 tuples)
If a clustered B+ tree index on $rname$, cost \approx 100 I/Os; if an unclustered index, up to 10,000 I/Os (even worst than entire file scan!)

Selection σ (Cont.)

- Important refinement for unclustered indexes
 1. Find qualifying data entries
 2. **Sort** the rid's of the data records to be retrieved by their **page-id** component
 3. Fetch rids in order. This ensures that each data page is looked at just once (though # of such pages likely to be higher than with clustering)

Cost of retrieving tuples = # of pages containing qualifying tuples

Selection σ (Cont.)

- General selection conditions
 - A Boolean combination (\wedge, \vee) of terms
 - Terms have the form:

$$attr \text{ op } constant, \text{ or, } attr1 \text{ op } attr2$$
 - Conditions expressed in conjunctive normal form (CNF)
 - A condition is a collection of **conjuncts** that are connected by \wedge
 - A conjunct contains consist of one or more **terms** connected by \vee
 - A conjunct containing \vee is said to be **disjunctive** (contain disjunction)

Selection σ (Cont.)

○ Selections without disjunction

Approach 1: find the most selective access path, retrieve tuples using it, and apply any remaining terms that don't match the access path

- **Most selective access path**: an index or file scan that is estimated to require the fewest page I/Os
- Terms that match the access path reduce # of tuples retrieved; other terms used to discard some retrieved tuples
- *day < 8/9/03 AND bid = 5 AND sid = 3*

A B+ tree index on *day* can be used; then, *bid = 5* and *sid = 3* must be checked for each retrieved tuple. Similarly, a hash index on *<bid, sid>* could be used; *day < 8/9/03* must then be checked

Selection σ (Cont.)

Approach 2 (if we have 2 or more matching indexes that use Alternatives (2) or (3) for data entries)

- Retrieve the rids of data records using each matching index, then **intersect** these sets of rids (we'll discuss intersection soon)
- Retrieve the records and apply any remaining terms
- *day < 8/9/03 AND bid = 5 AND sid = 3*

If we have a B+ tree index on *day* and an index on *sid*, both using Alternative (2), we can retrieve rids of records satisfying *day < 8/9/03* using the first, rids of records satisfying *sid = 3* using the second, intersect, retrieve records and check *bid = 5*

* We don't discuss selections with disjunction

Projection π

$$\pi_{attr1, attr2, \dots, attrn}(R)$$

○ Implementing projection

- Remove unwanted attributes
- Eliminate duplicates (based on sorting / hashing)

```
SELECT DISTINCT R.sid, R.bid
FROM Reserves R
```

Projection π (Cont.)


○ Projection based on sorting

- Modify pass 0 of external merge sort to eliminate unwanted fields
 - Runs of about 2B pages are produced, but tuples in runs are smaller than input tuples (size ratio depends on the total size of dropped fields)
- Modify merging passes to eliminate duplicates
 - Less result tuples than input (difference depends on # of duplicates)
- I/O cost
 - In Pass 0, read the original relation (*M* pages), write out the same number of smaller tuples
 - In merging passes, fewer pages are written out in each pass

Projection π (Cont.)

○ Projection based on hashing

- Used when # of buffer pages (B) is much larger than # of pages in relation R
- Cost: read R for partitioning, write out each tuple with fewer fields (therefore fewer pages); they are read in the next phase, in-memory hash table is written out for each partition



Projection π (Cont.)

○ Phase 1: partitioning

Read R using one input buffer; for each tuple, discard unwanted fields, apply hash function $h1$ to choose one of $B-1$ output buffers

- Result is $B-1$ partitions (with no unwanted fields)
- 2 tuples from different partitions must be distinct

○ Phase 2: duplicate elimination in each partition

For each partition: read it in, one page at a time; hash its tuples using $h2$ ($\neq h1$), then insert tuples into an in-memory hash table, discard duplicates; write hash table to the result file

- If a tuple hashes to the same value as an existing tuple, compare the two to check whether duplicates

Projection π (Cont.)

○ Sort-based approach is the standard

- Better handling of non-uniformly distributed values
- Result is sorted

○ If an index on the relation contains all wanted attributes in its search key, can do **index-only** scan

- Apply projection techniques to data entries (much smaller!)

○ If an **ordered** (i.e., tree) index contains all wanted attributes as **prefix** of search key, can do even better

- Retrieve data entries in order (index-only scan), discard unwanted fields, compare adjacent tuples to check for duplicates

Set Operations

○ $R \cap S$, $R \times S$, $R \cup S$, $R - S$

○ Intersection (\cap) and cross-product (\times) are implemented as special cases of join

- \cap : equality on all fields as the join condition
- \times : no join condition

○ Implementation of union (\cup) and set-difference ($-$) are similar